Foundations of Deep Learning

ALF

Alfredo Canziani

@alfcnz

Convolutional Neural Nets

Exploiting stationarity, locality, and compositionality of natural data

Signals can be represented as vectors

$$\boldsymbol{x} = [x_1 \;\; x_2 \;\; x_3 \;\; \ldots \;\; x_t \;\; \ldots]^\top$$

x_t are waveform heights

$$\boldsymbol{x} = [x_{11} \;\; x_{12} \;\; \ldots \;\; x_{1n} \;\; x_{21} \;\; x_{22} \;\; \ldots]^\top$$

x_{ij} are pixel values

"John picked up the apple"

$$\boldsymbol{x} = [x_1 \;\; x_2 \;\; x_3 \;\; x_4 \;\; x_5]^\top$$

x_t are one-hot vectors

Signals can be represented as vectors

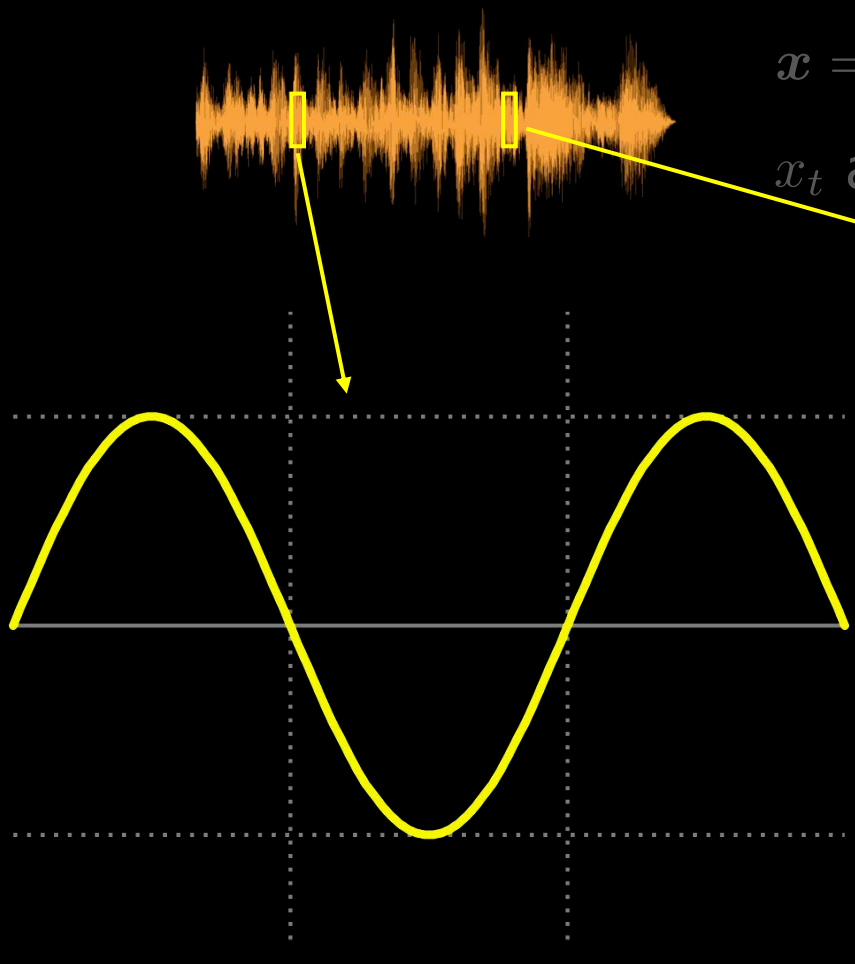

"John picked up the apple"

$\boldsymbol{x} = [$

x_t ar

x_1

e

$[x_1 \ x_2 \ x_3 \ x_4 \ x_5]^\top$

x_t are one-hot vectors

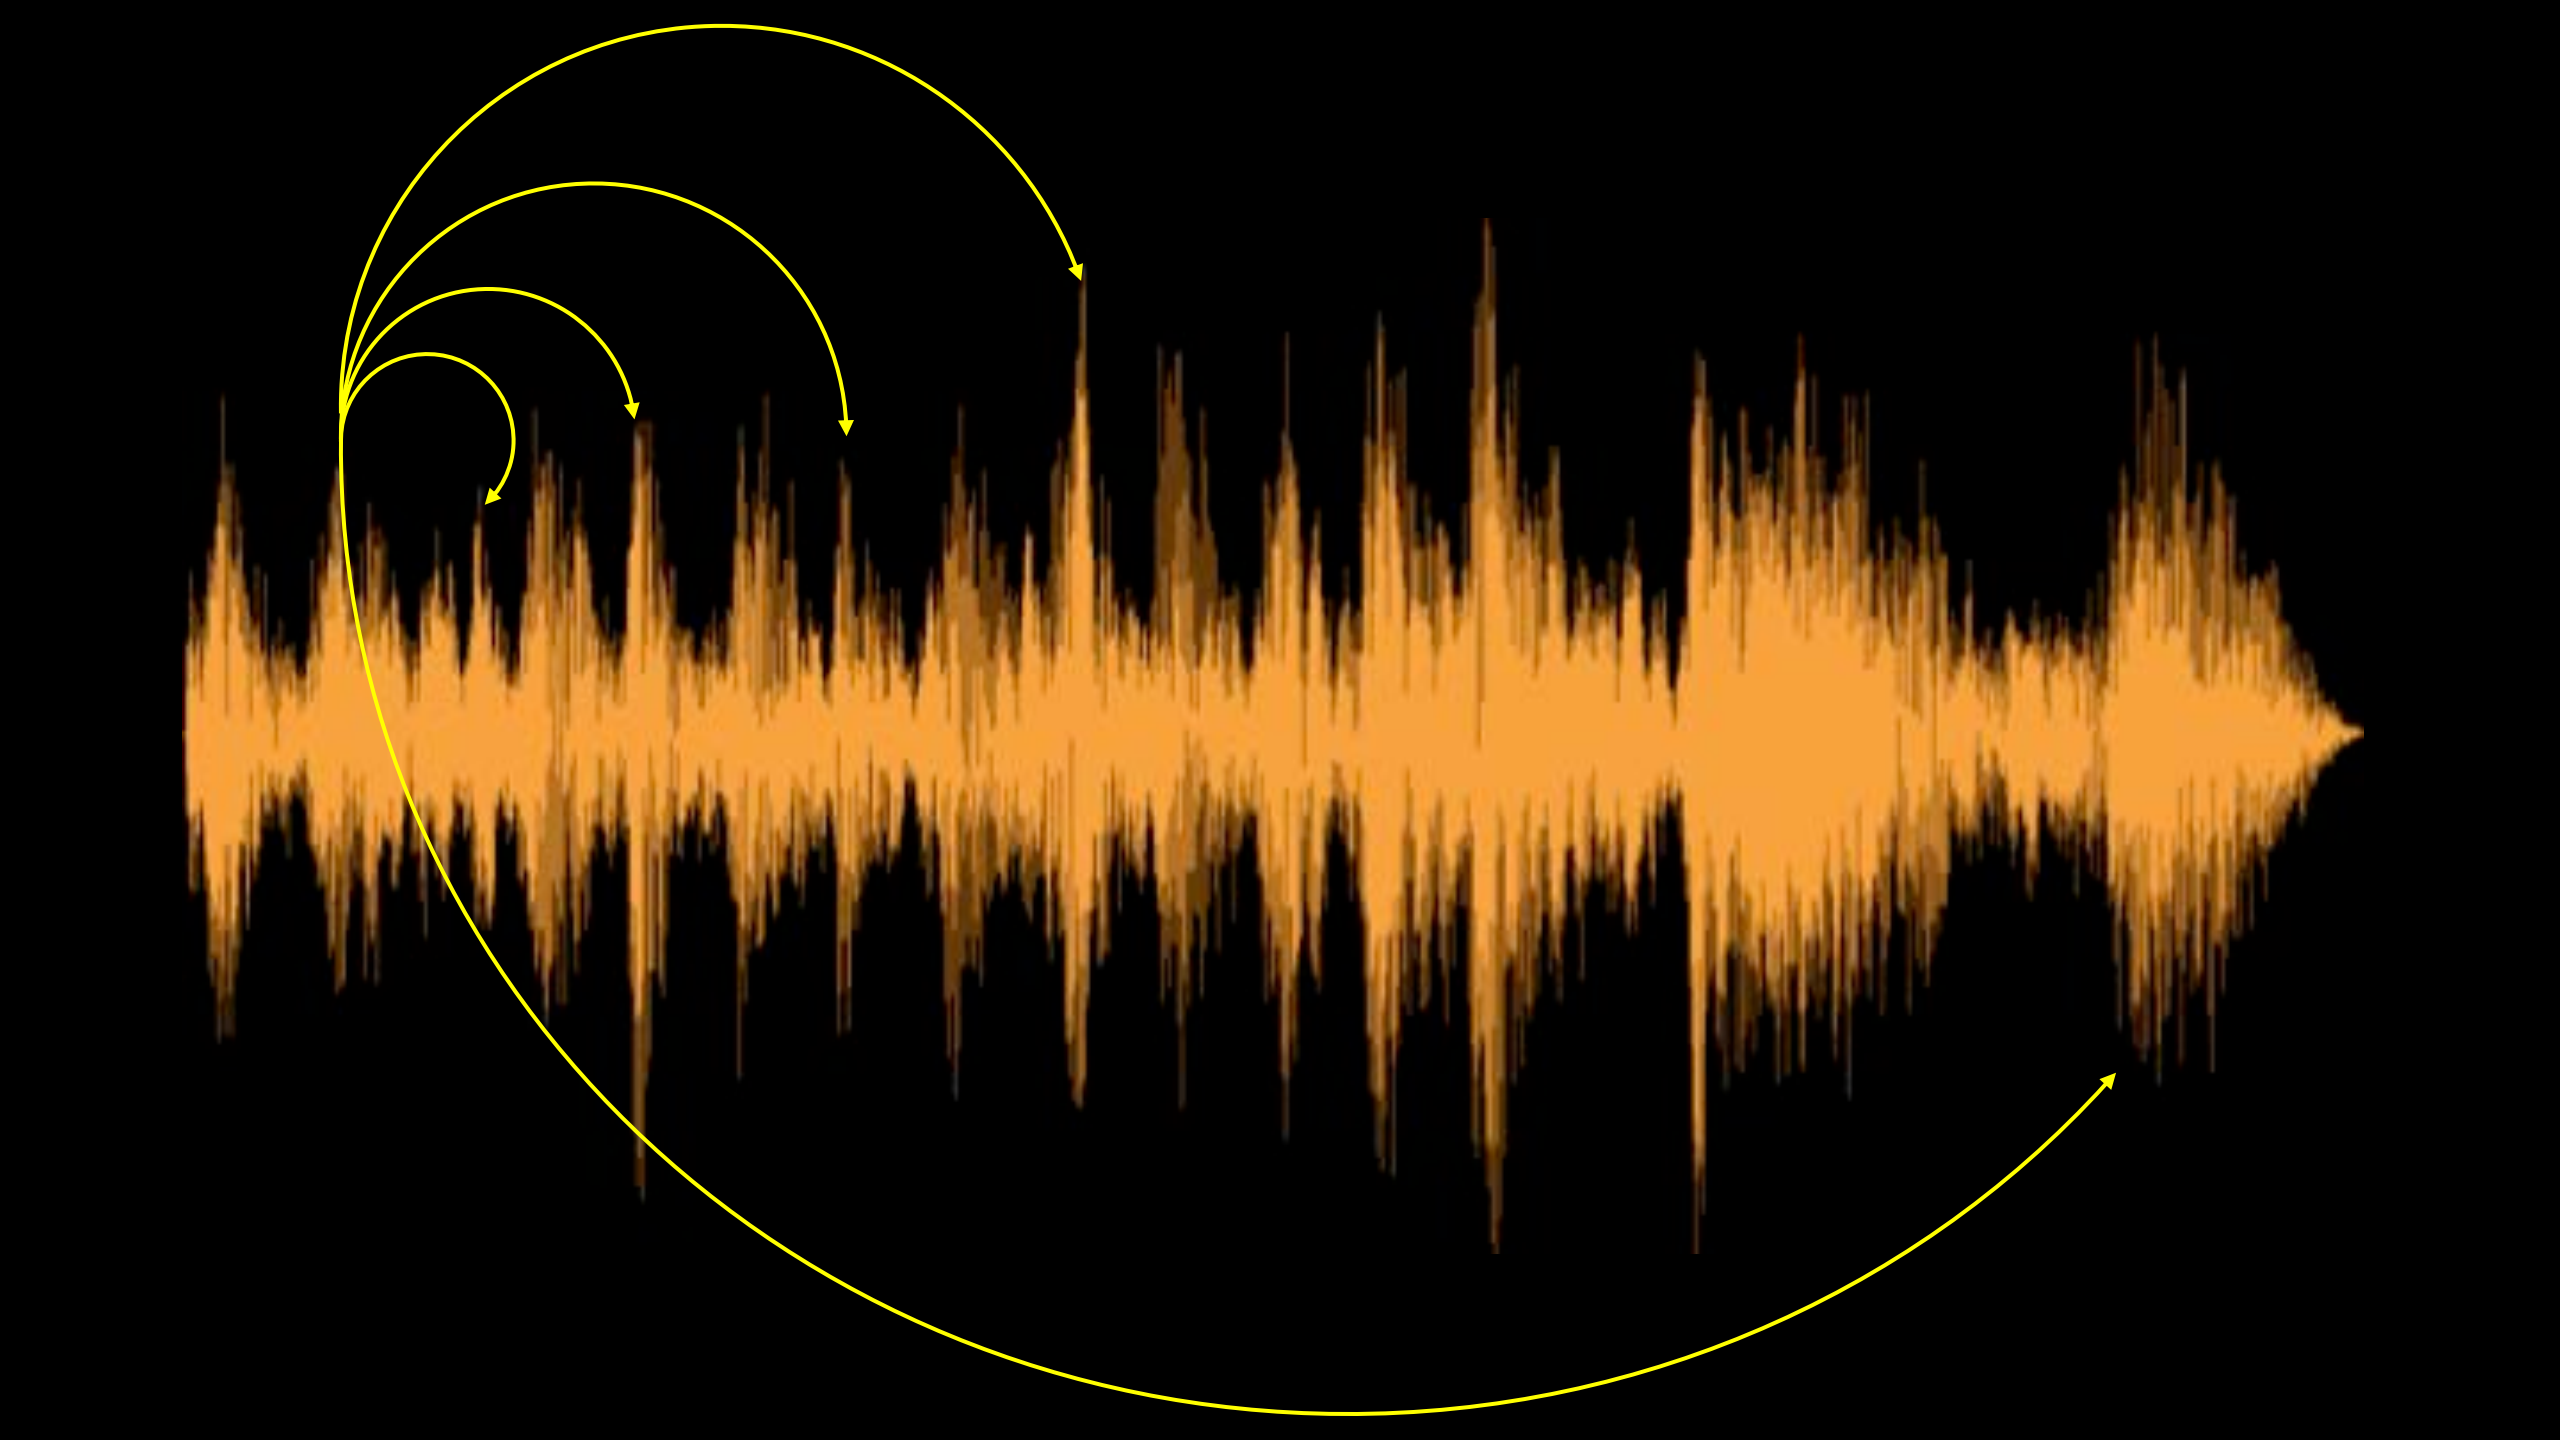

Signals can be represented as vectors

$$\boldsymbol{x} = [x_1 \;\; x_2 \;\; x_3 \;\; \ldots \;\; x_t \;\; \ldots]^\top$$

x_t are waveform heights

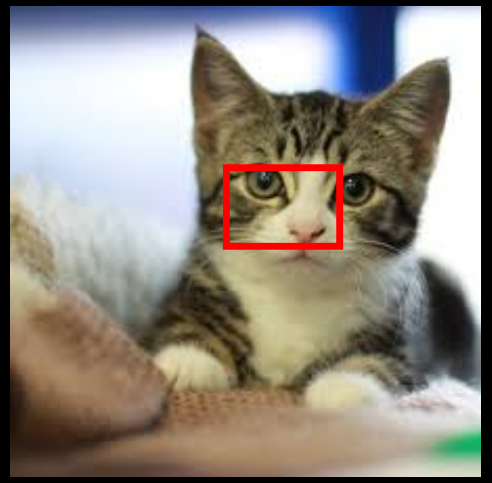

$$\boldsymbol{x} = [x_{11} \;\; x_{12} \;\; \ldots \;\; x_{1n} \;\; x_{21} \;\; x_{22} \;\; \ldots]^\top$$

x_{ij} are pixel values

"John picked up the apple"

$$\boldsymbol{x} = [x_1 \;\; x_2 \;\; x_3 \;\; x_4 \;\; x_5]^\top$$

x_t are one-hot vectors

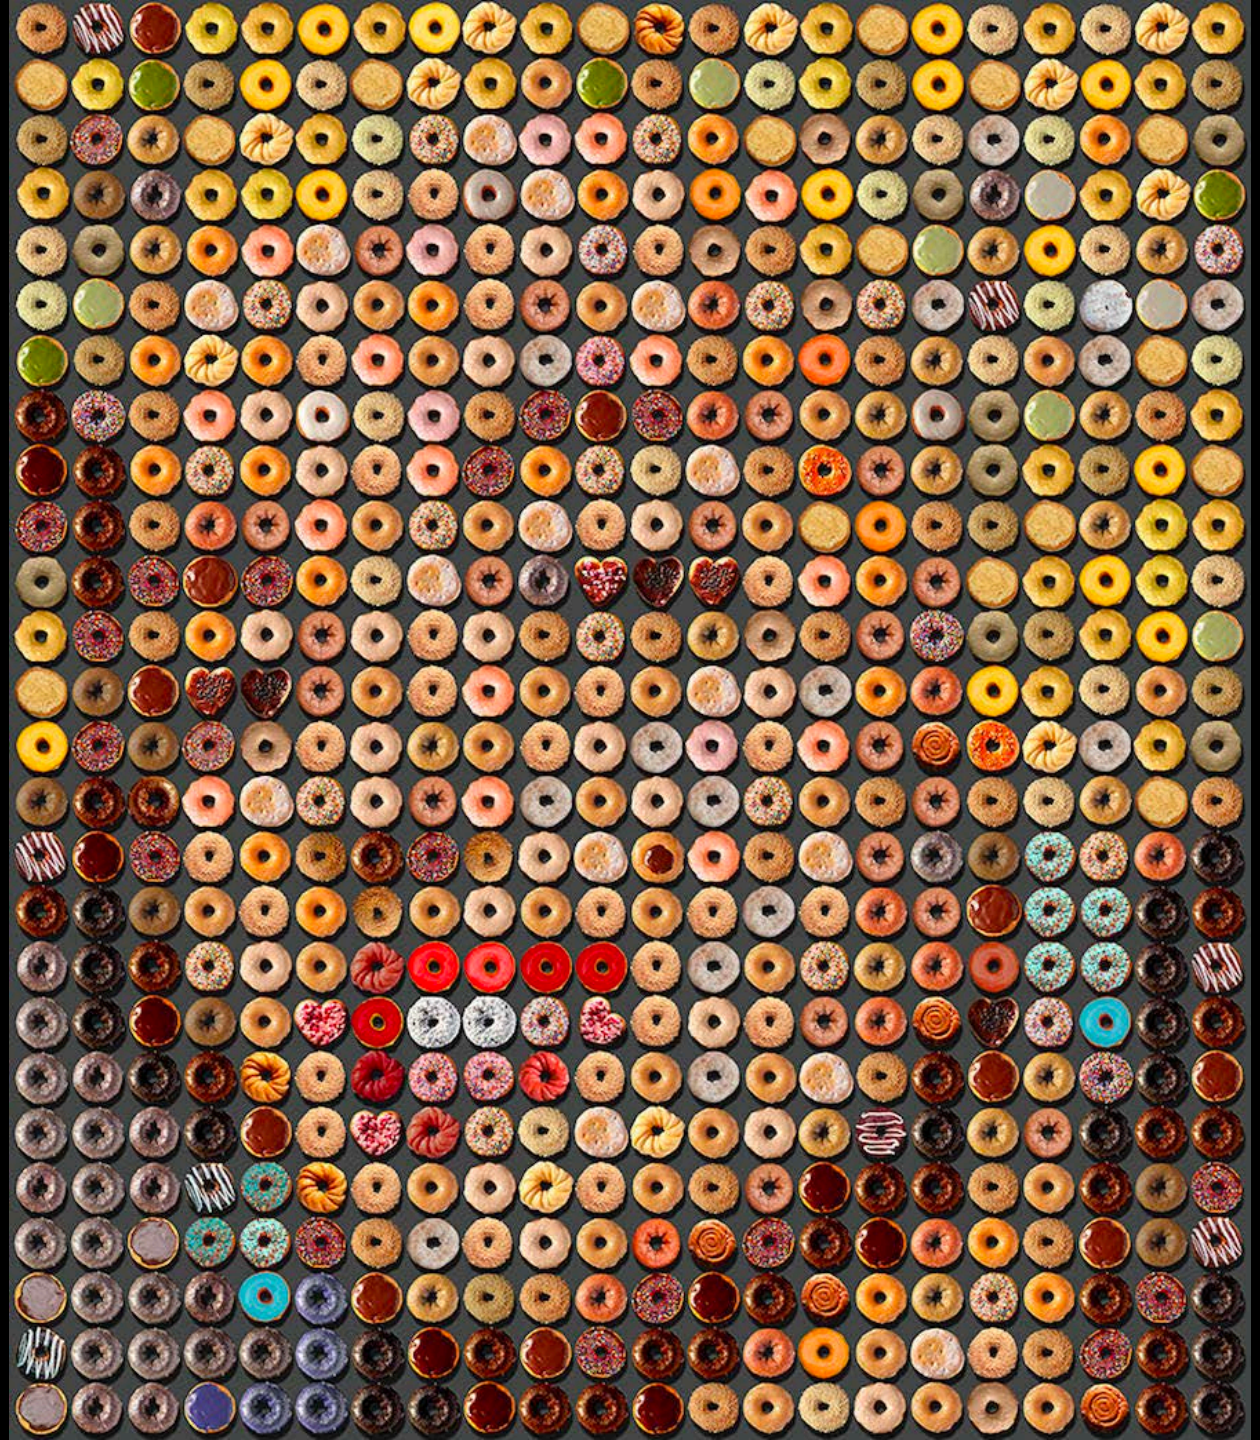

Fully connected (FC) layer

j-th row of $W^{(1)}$

$$a_j^{(2)} = f(\boldsymbol{w}^{(j)}\boldsymbol{x} + b_j) = f\left(\left(\sum_{i=1}^{n} w_i^{(j)} x_i\right) + b_j\right)$$

$$\boldsymbol{h} = f(\boldsymbol{W_h}\boldsymbol{x} + \boldsymbol{b_h})$$

$$\hat{\boldsymbol{y}} = g(\boldsymbol{W_y}\boldsymbol{h} + \boldsymbol{b_y})$$

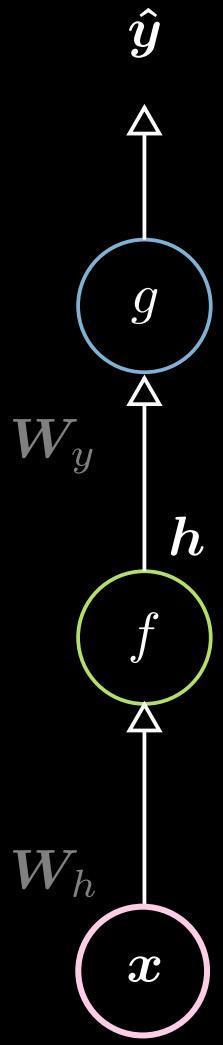

$f, g = (\cdot)^+, \sigma(\cdot),$
$\tanh(\cdot), \text{soft}(arg)\max(\cdot)$

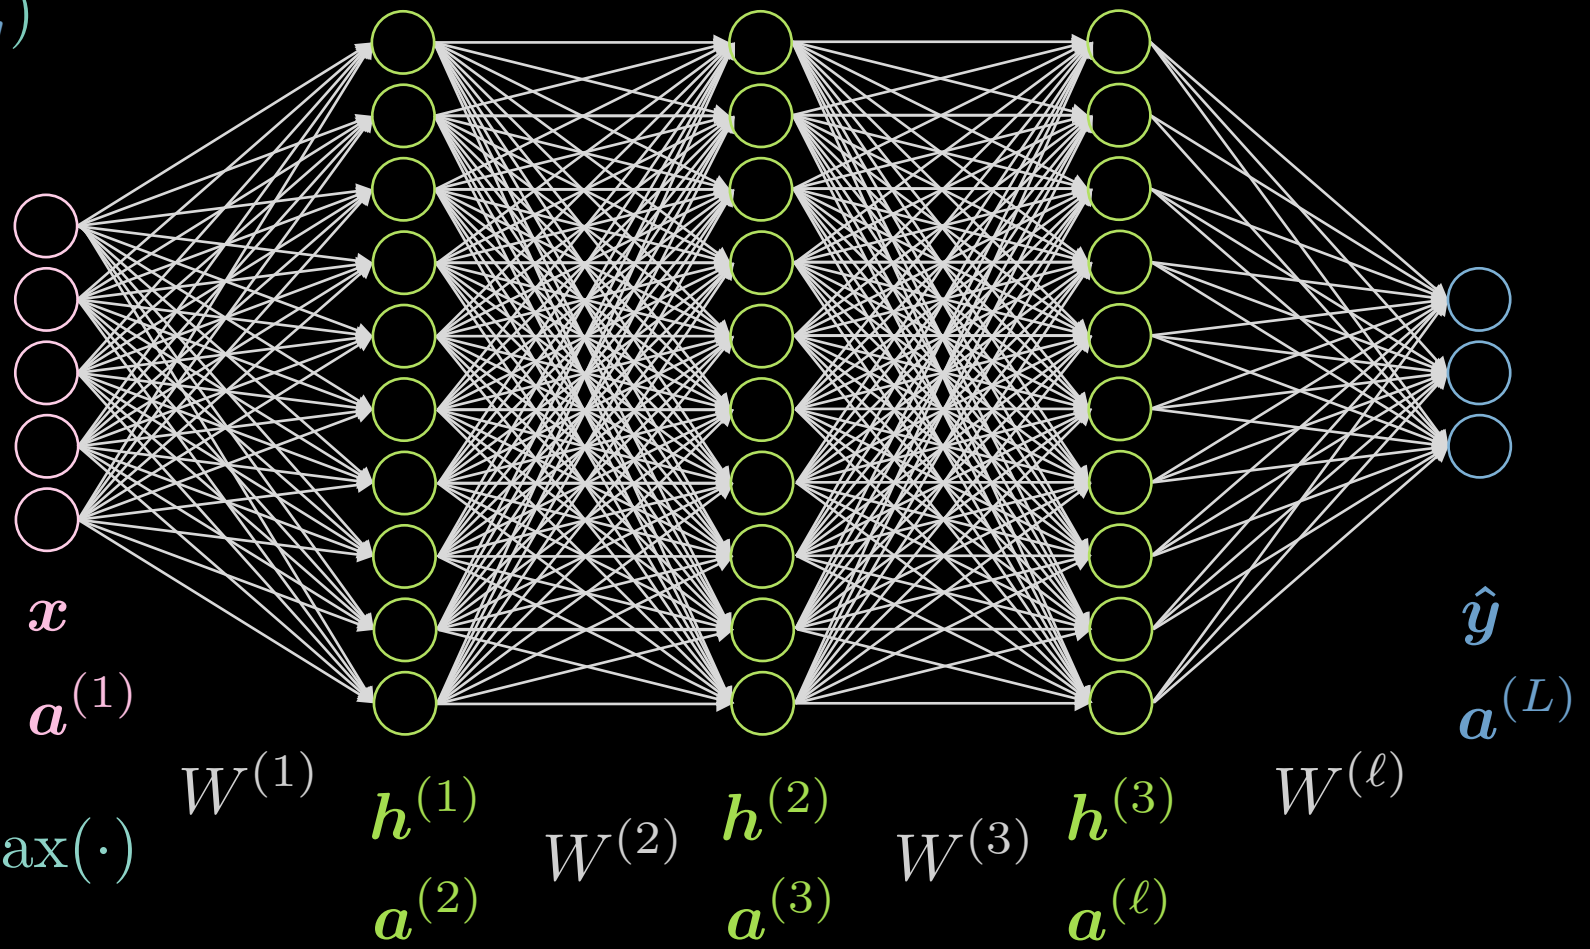

RF: receptive field

Locality ⇒ sparsity

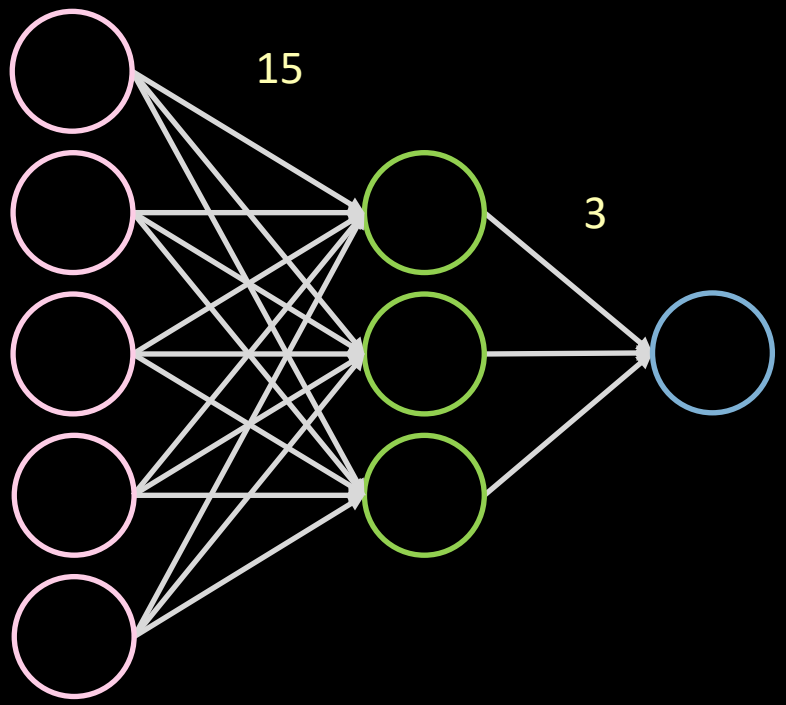

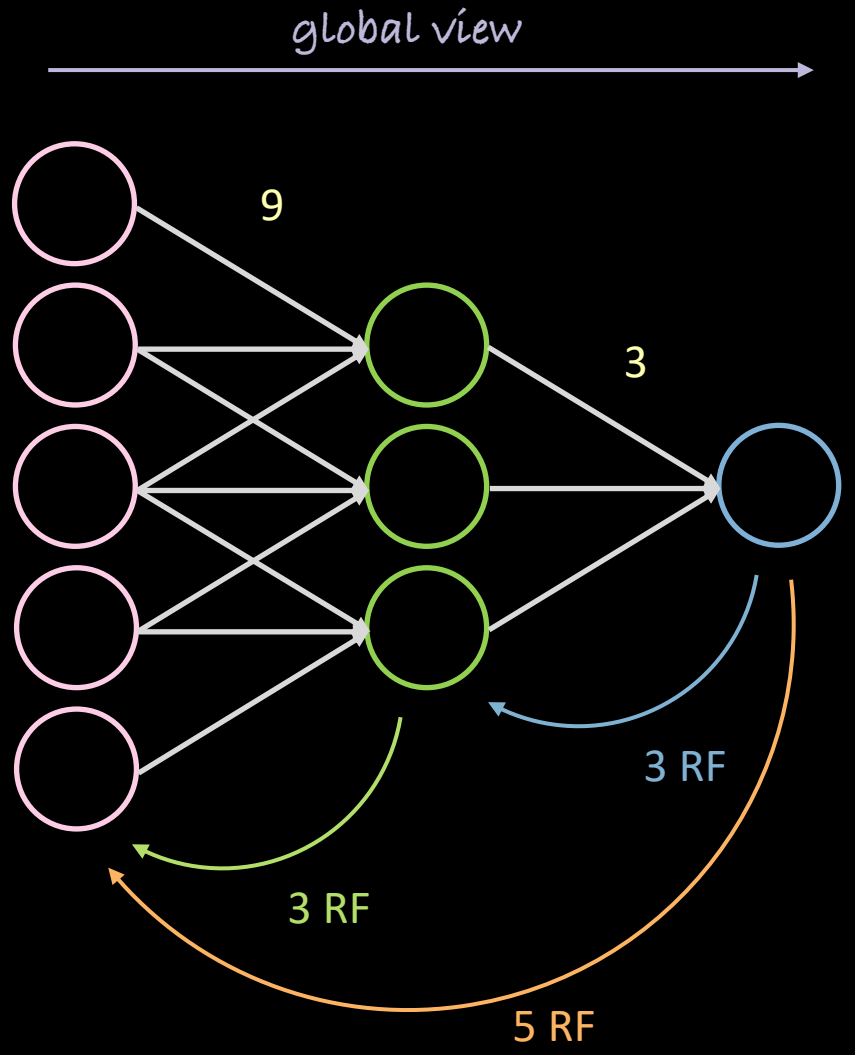

Stationarity ⇒ parameters sharing

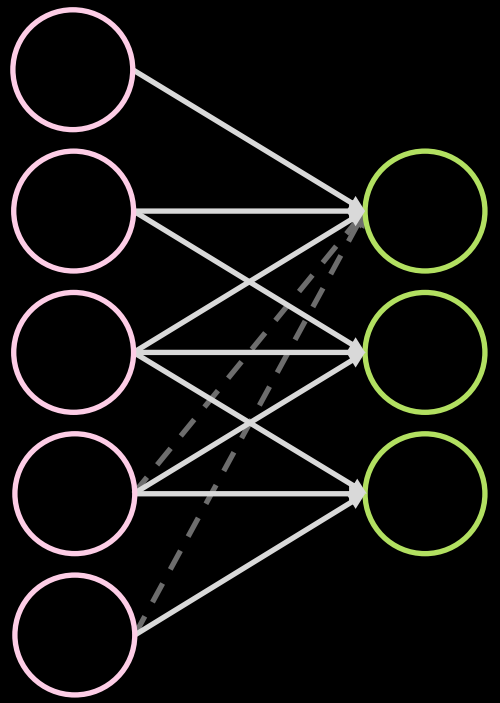

Parameters sharing

- faster convergence
- better generalisation
- not constrained to input size
- kernel independence
  ⇒ high parallelisation

Connection sparsity

- reduced amount of computation

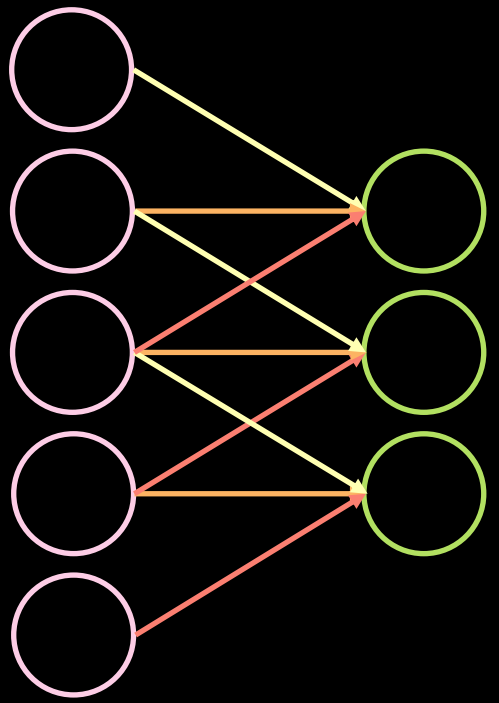

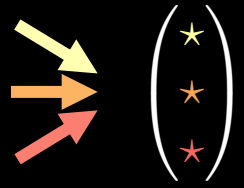

Generalisation error reduction

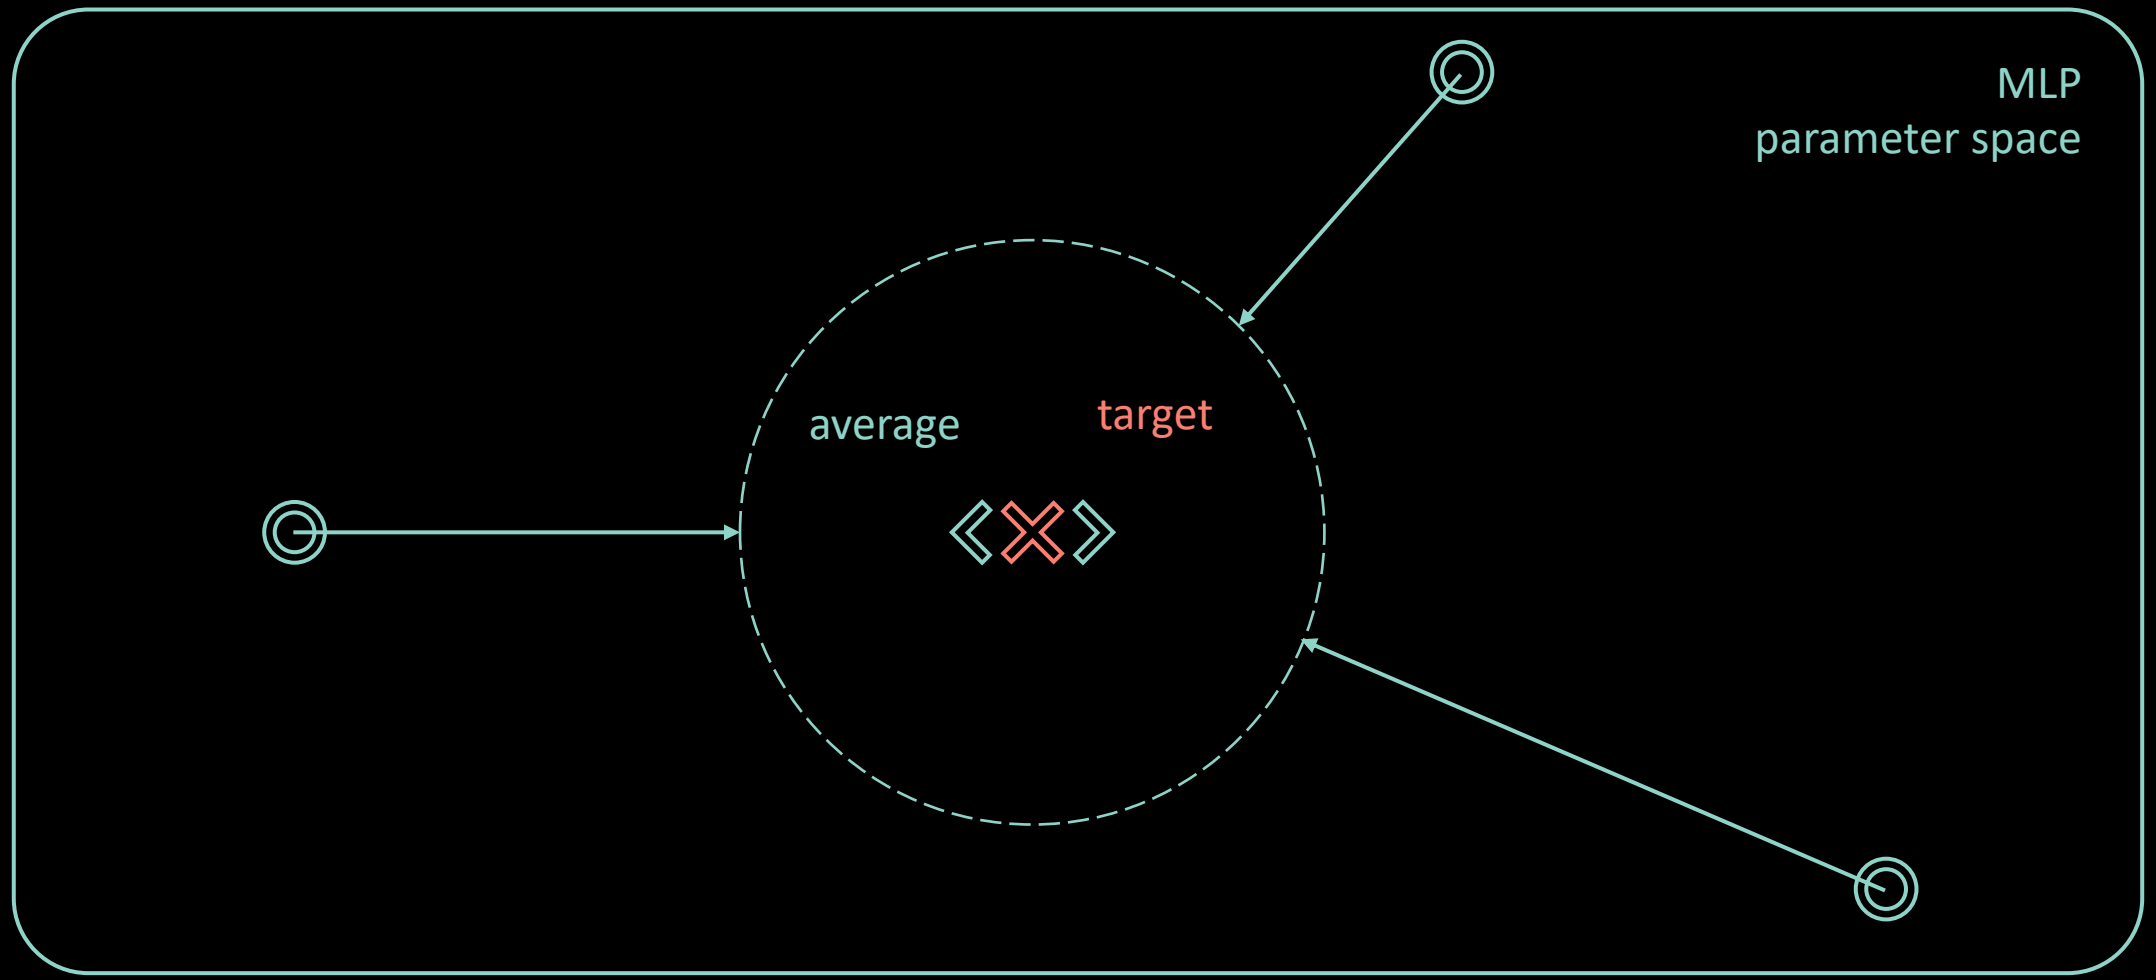

Generalisation error reduction

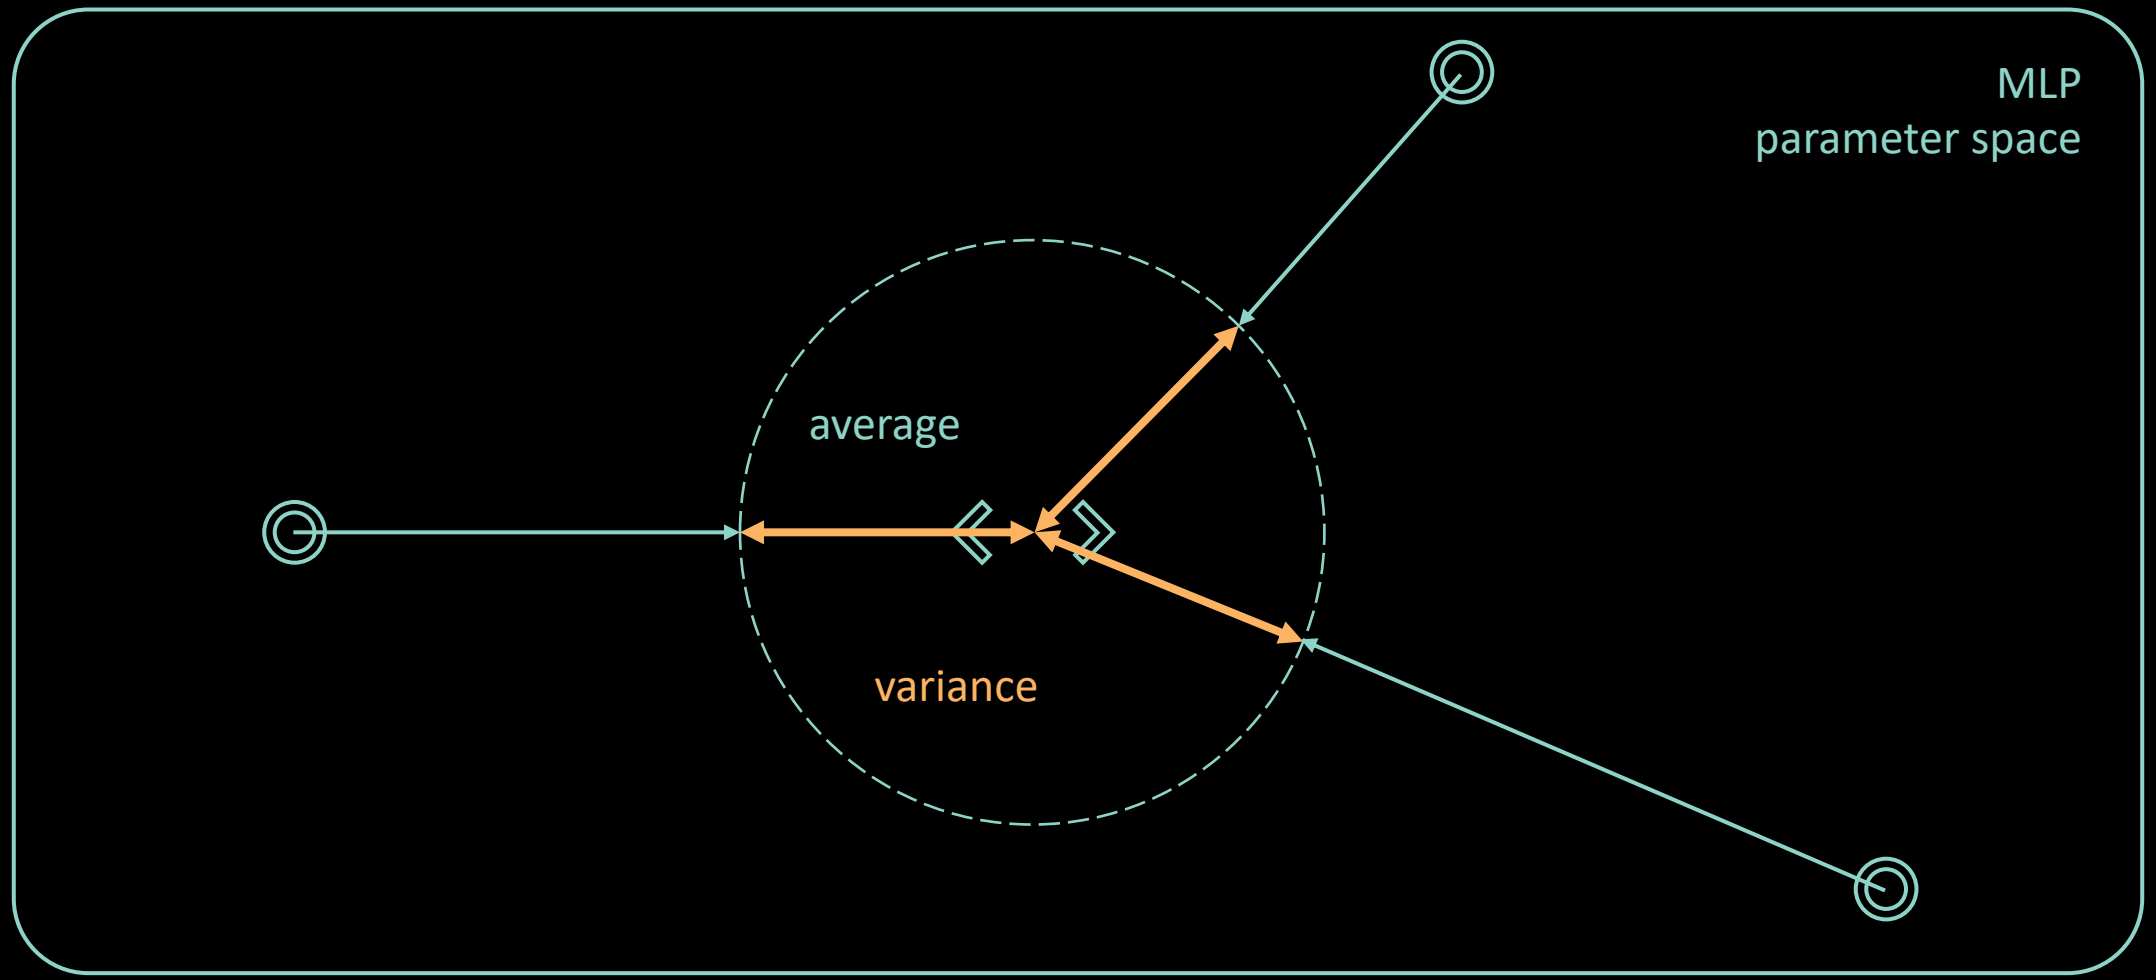

Generalisation error reduction

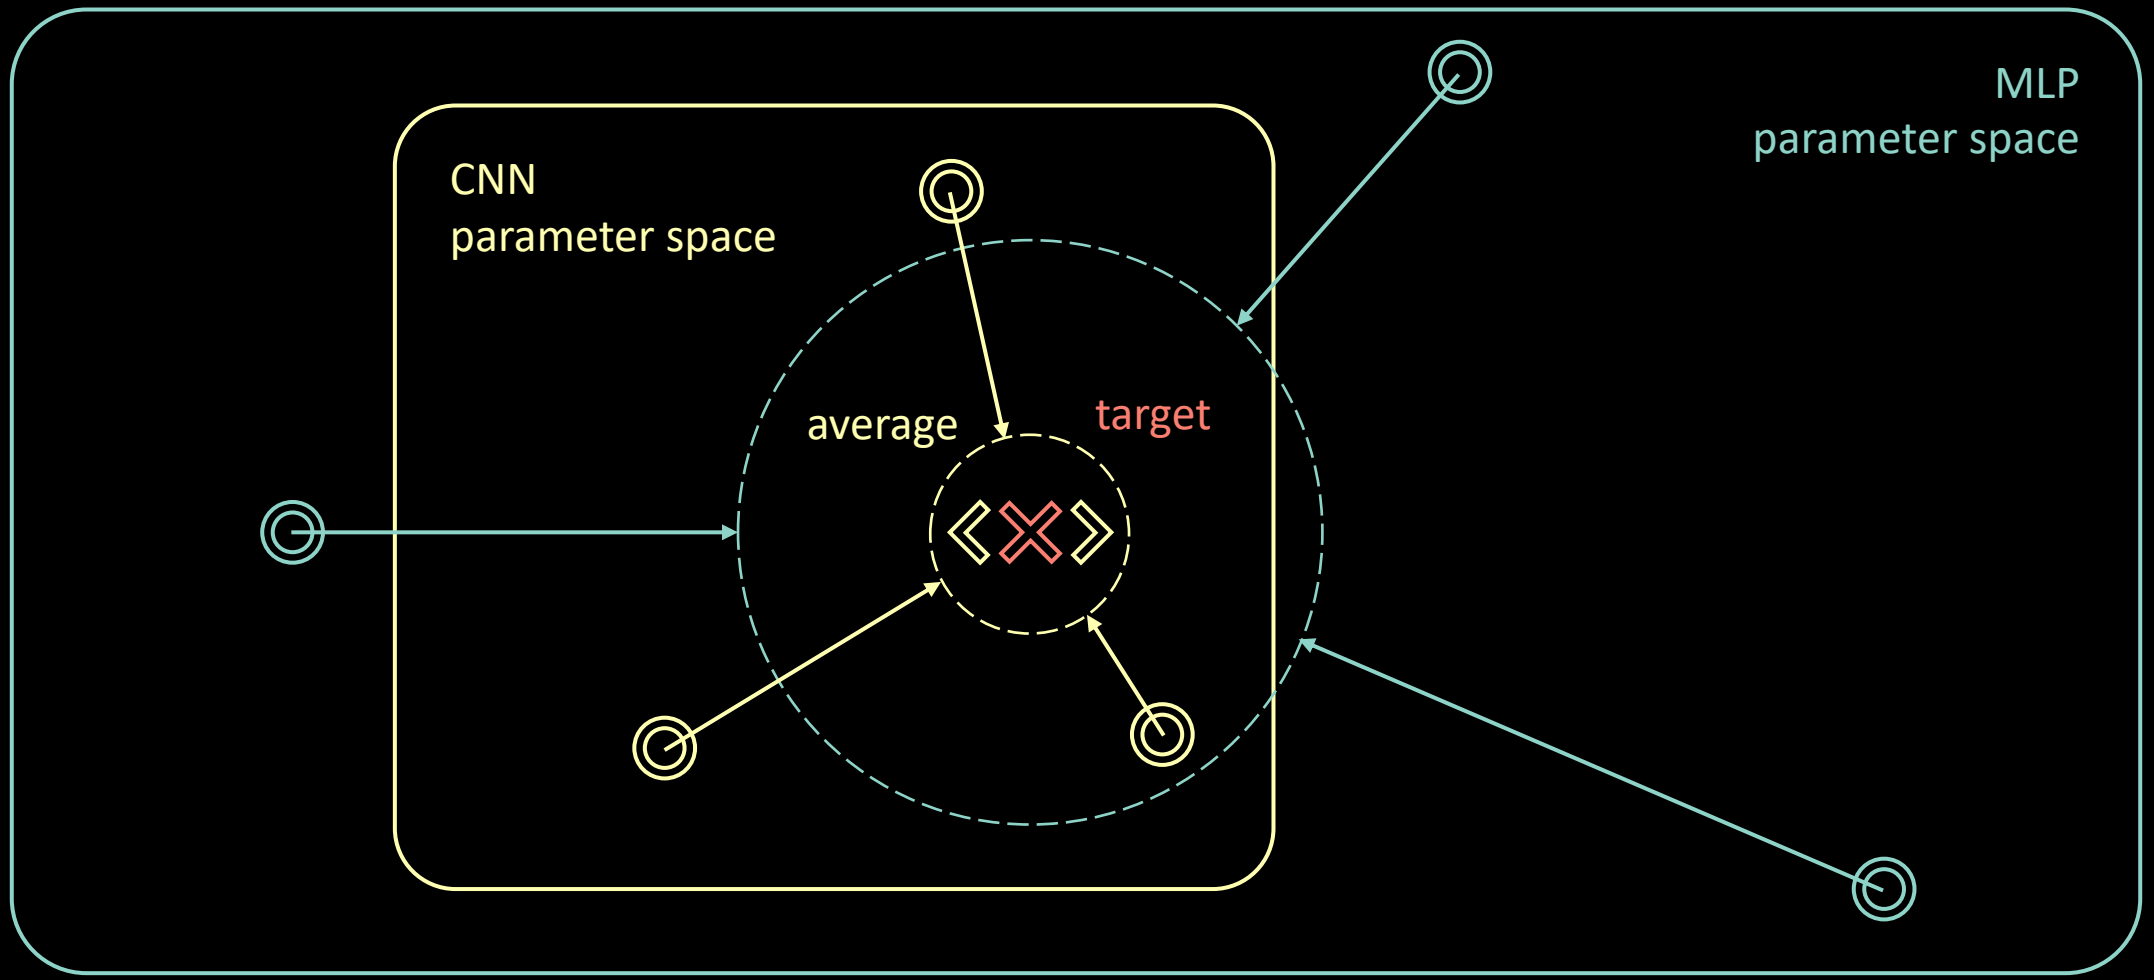

Generalisation error reduction

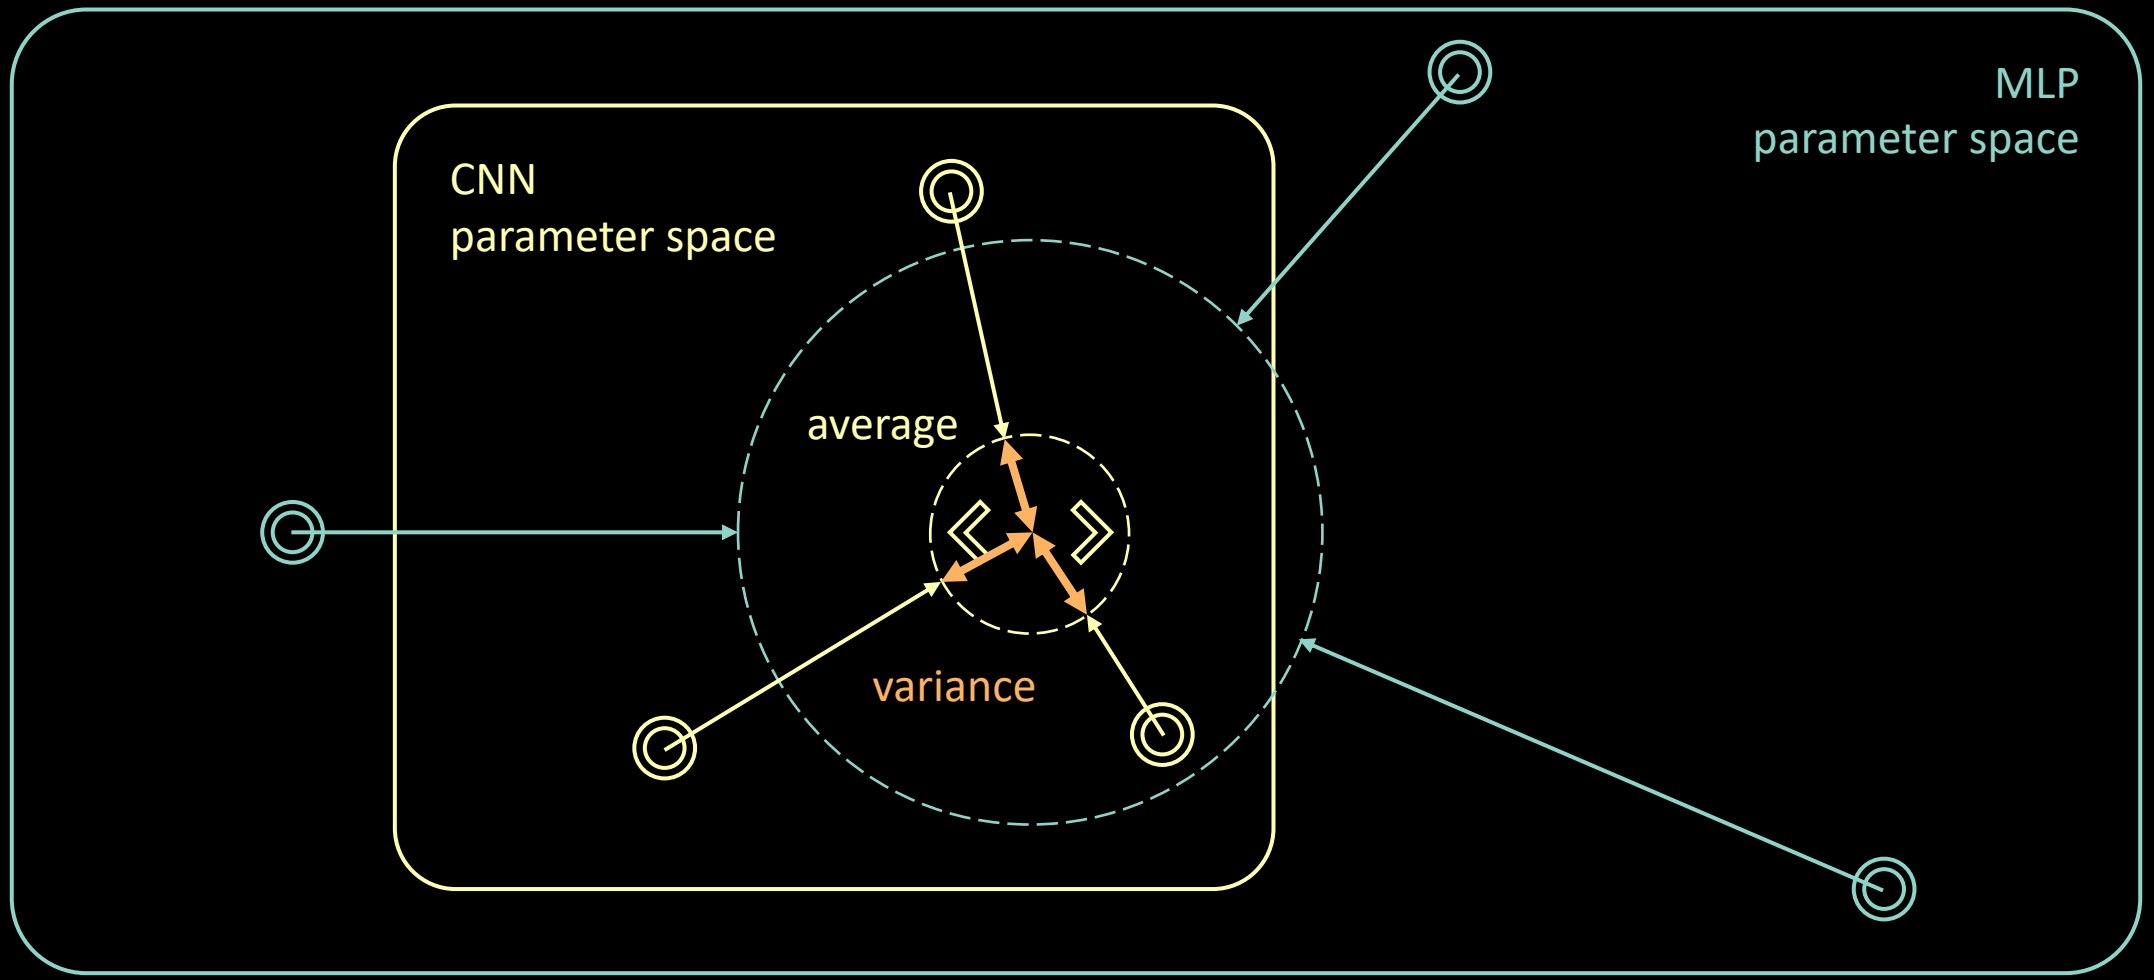

Generalisation error reduction

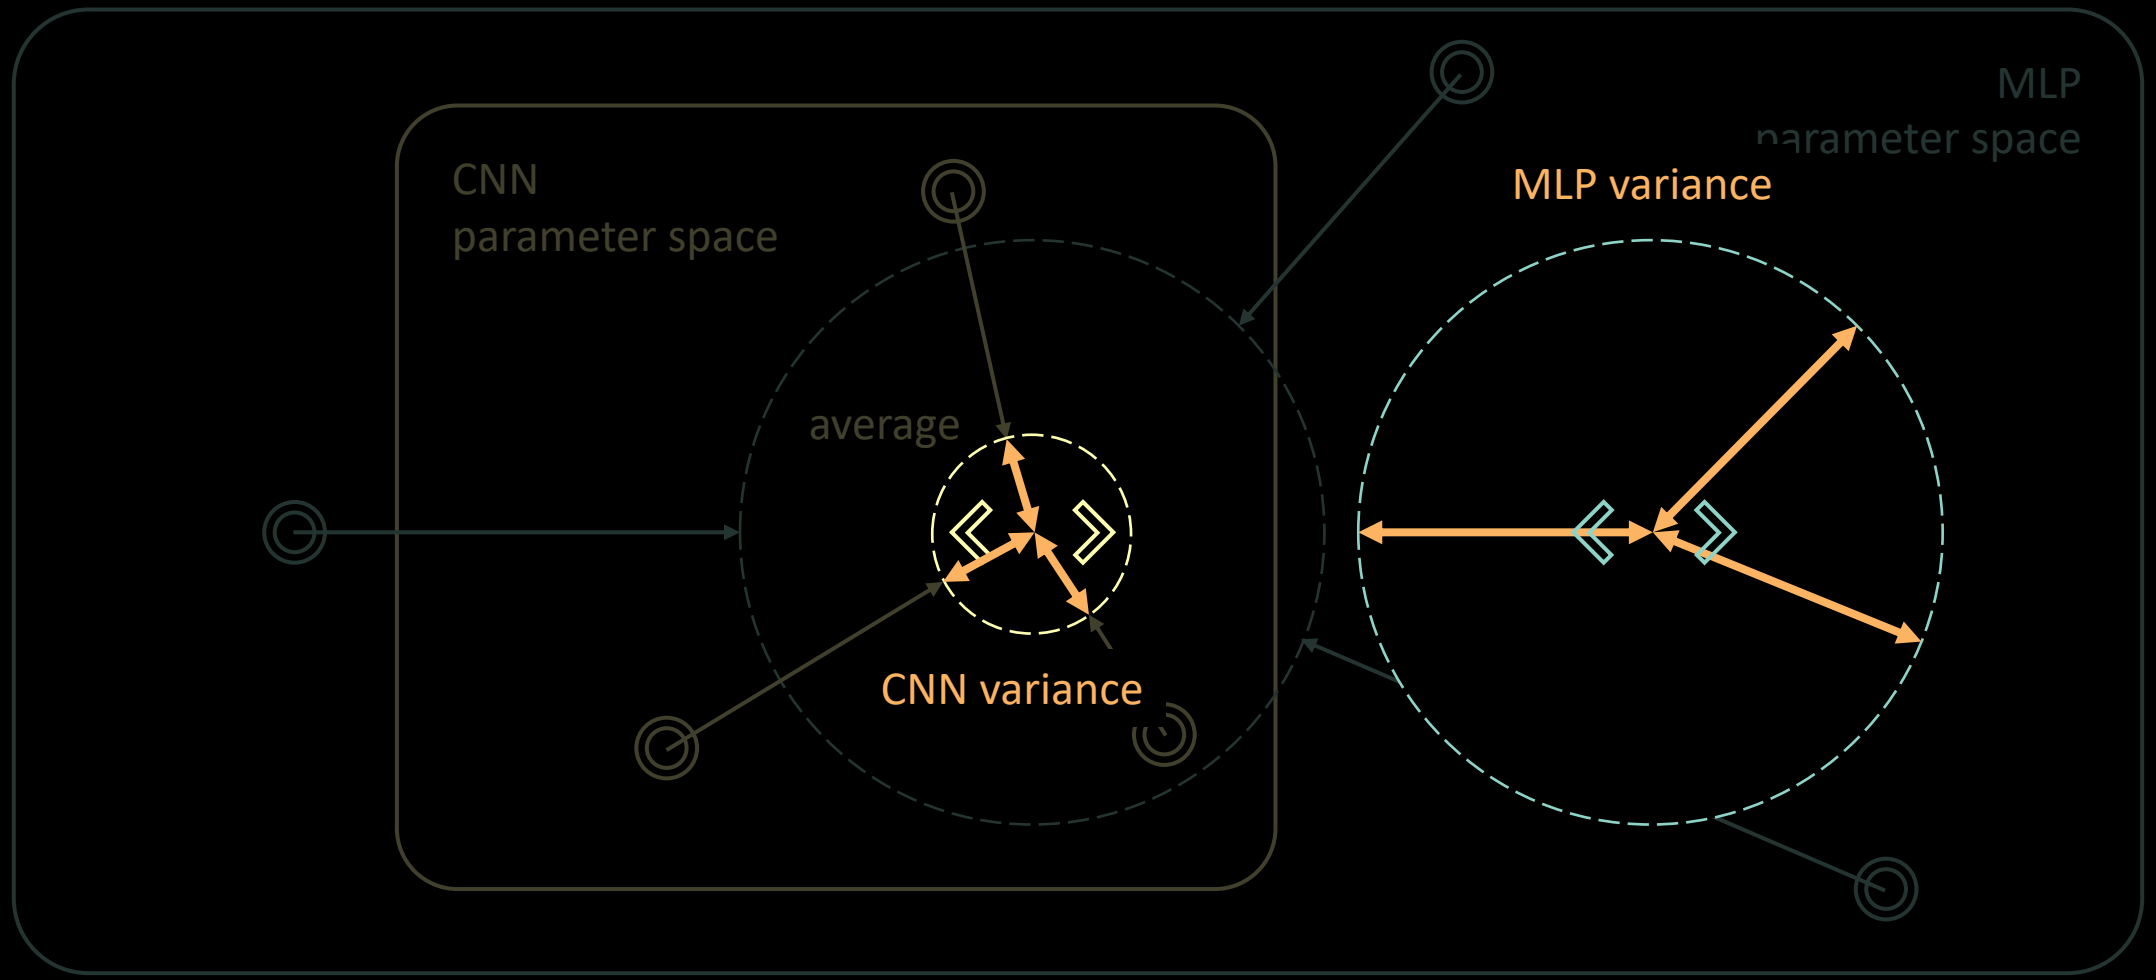

Misspecification of model constraints

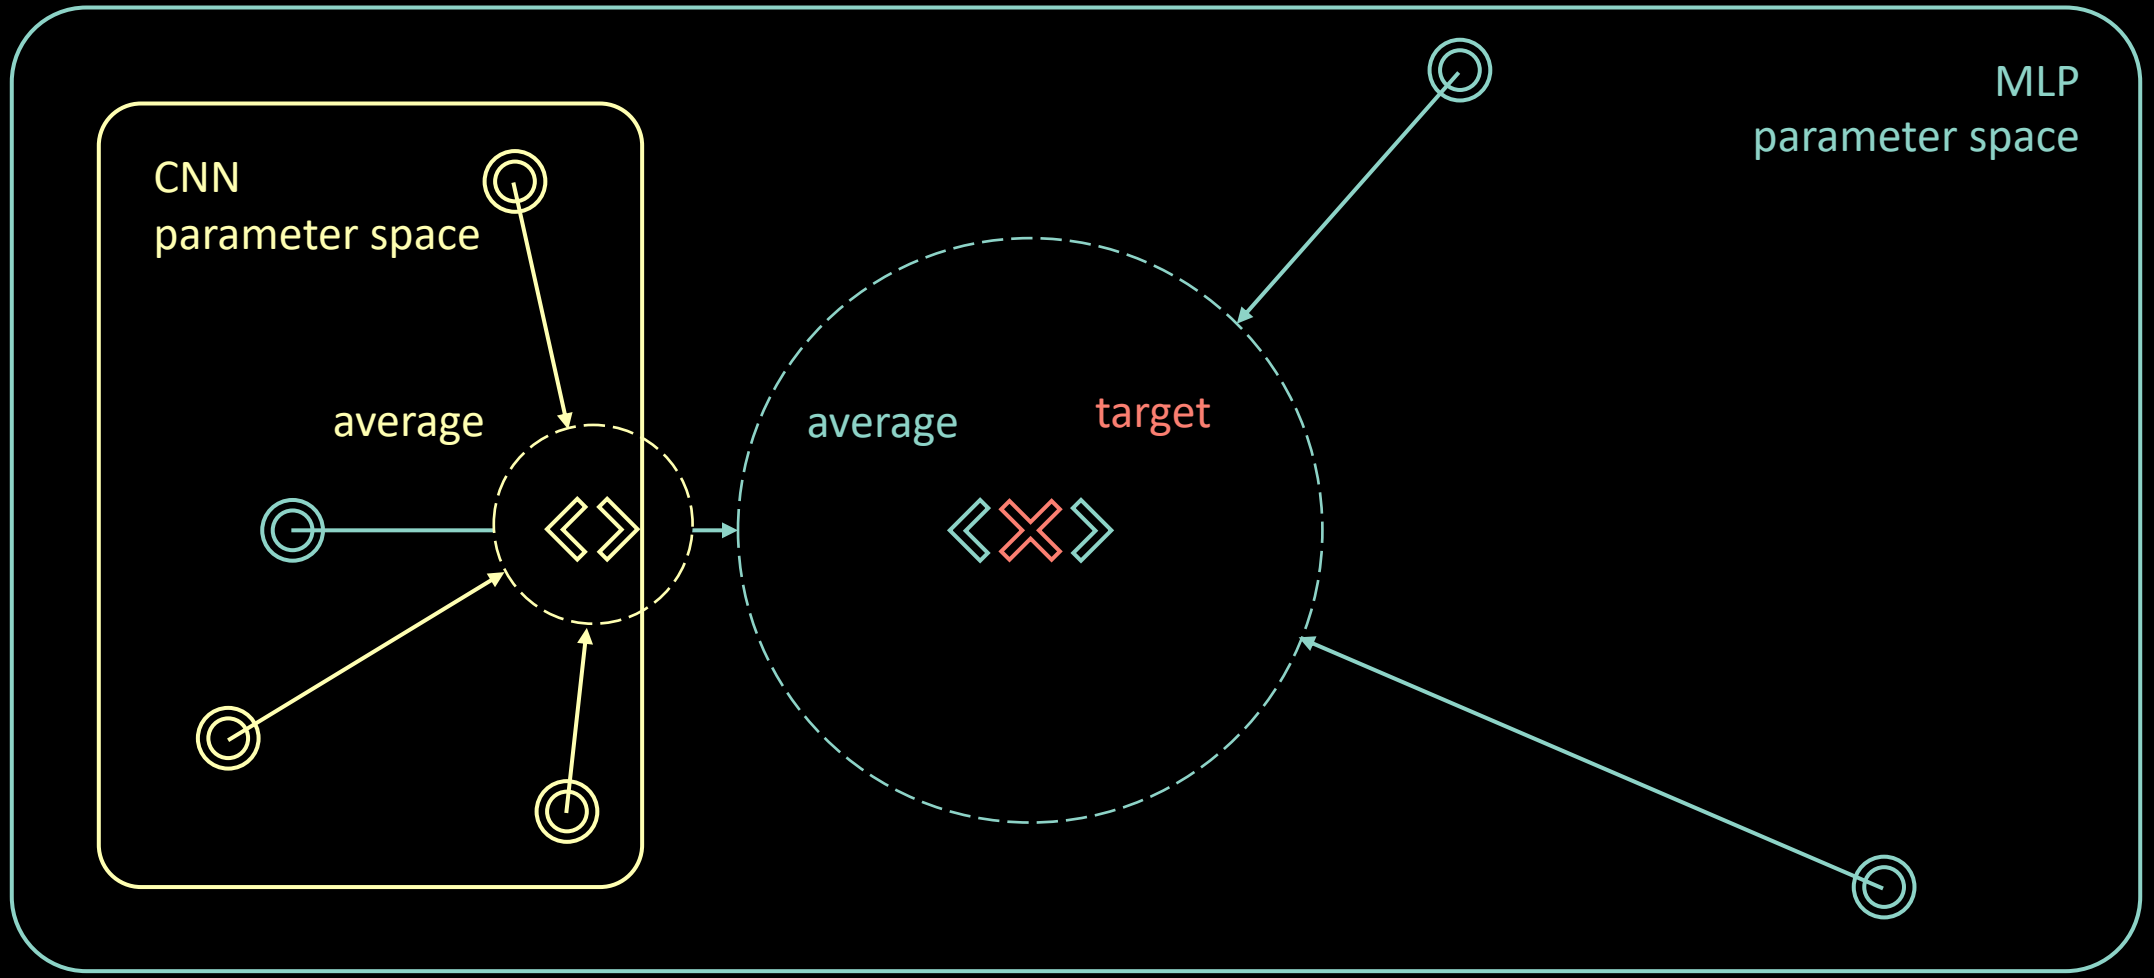

Misspecification of model constraints

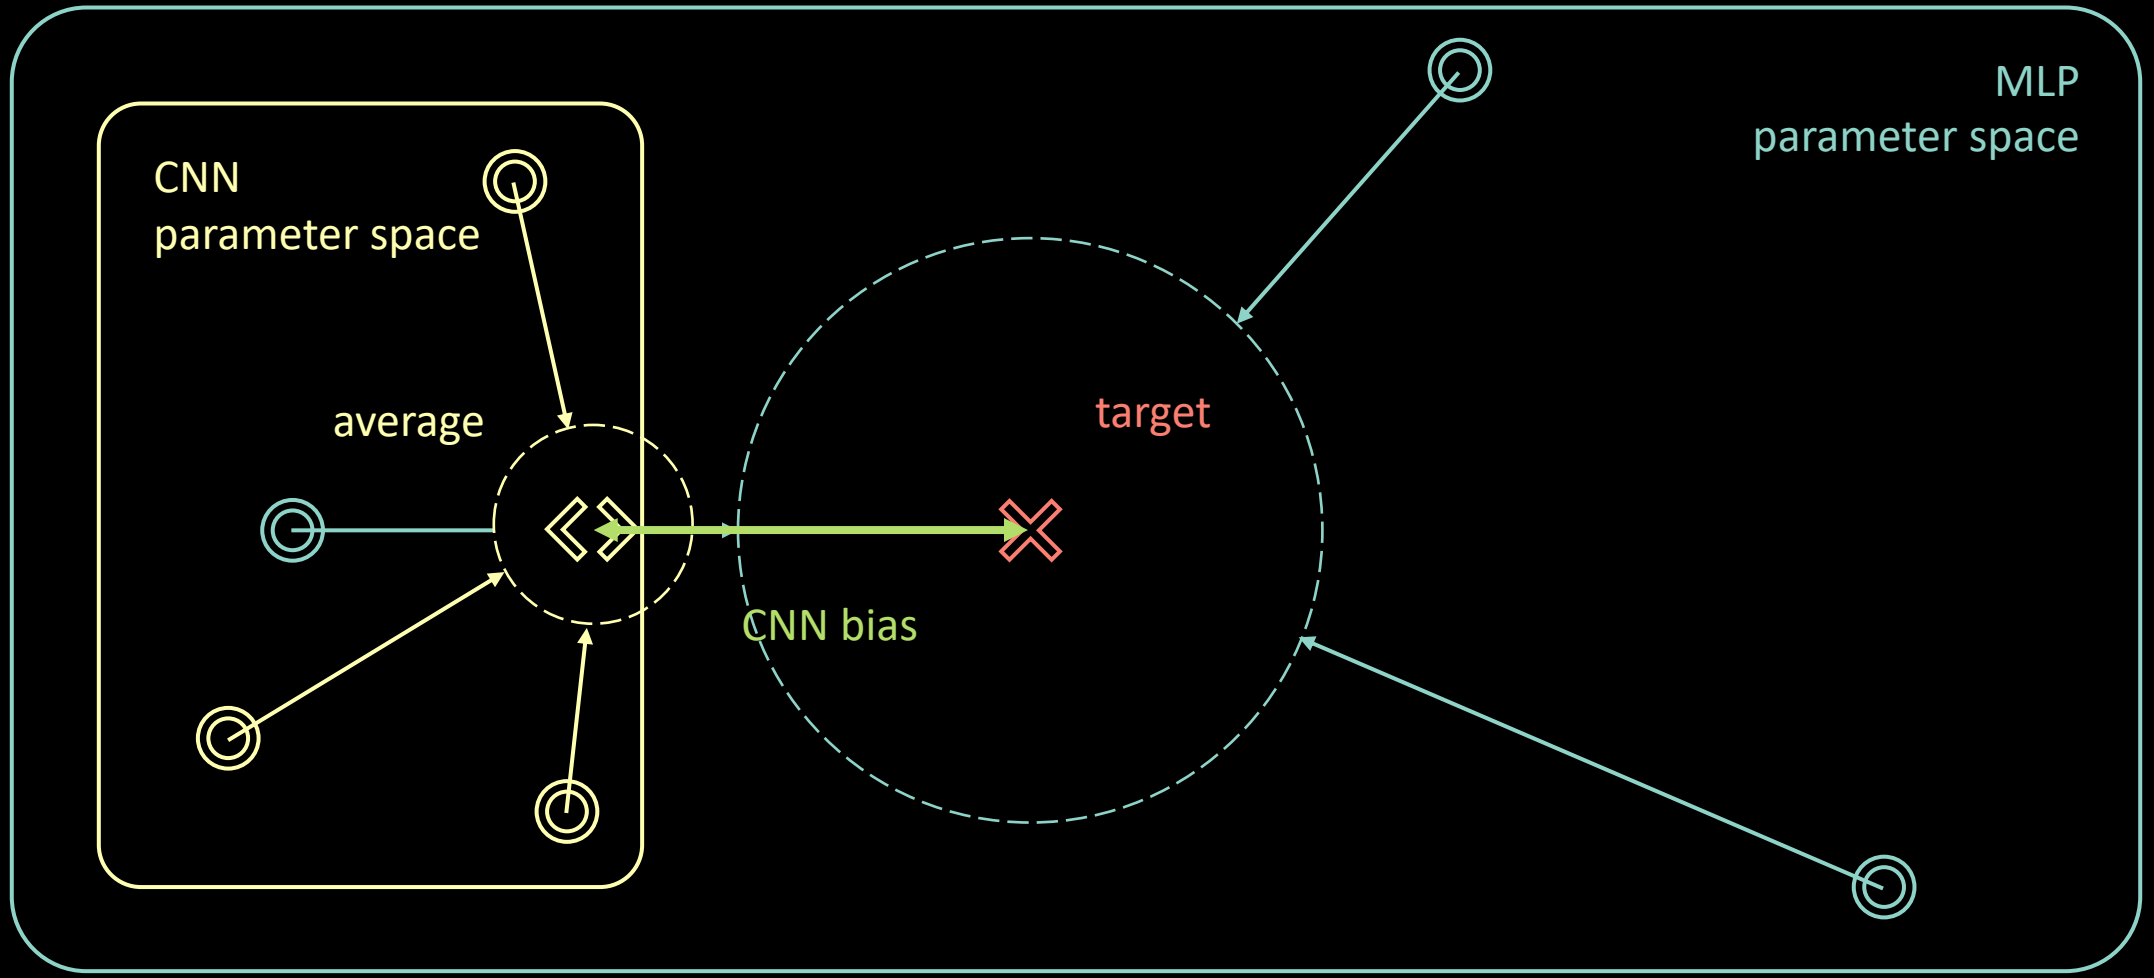

Misspecification of model constraints

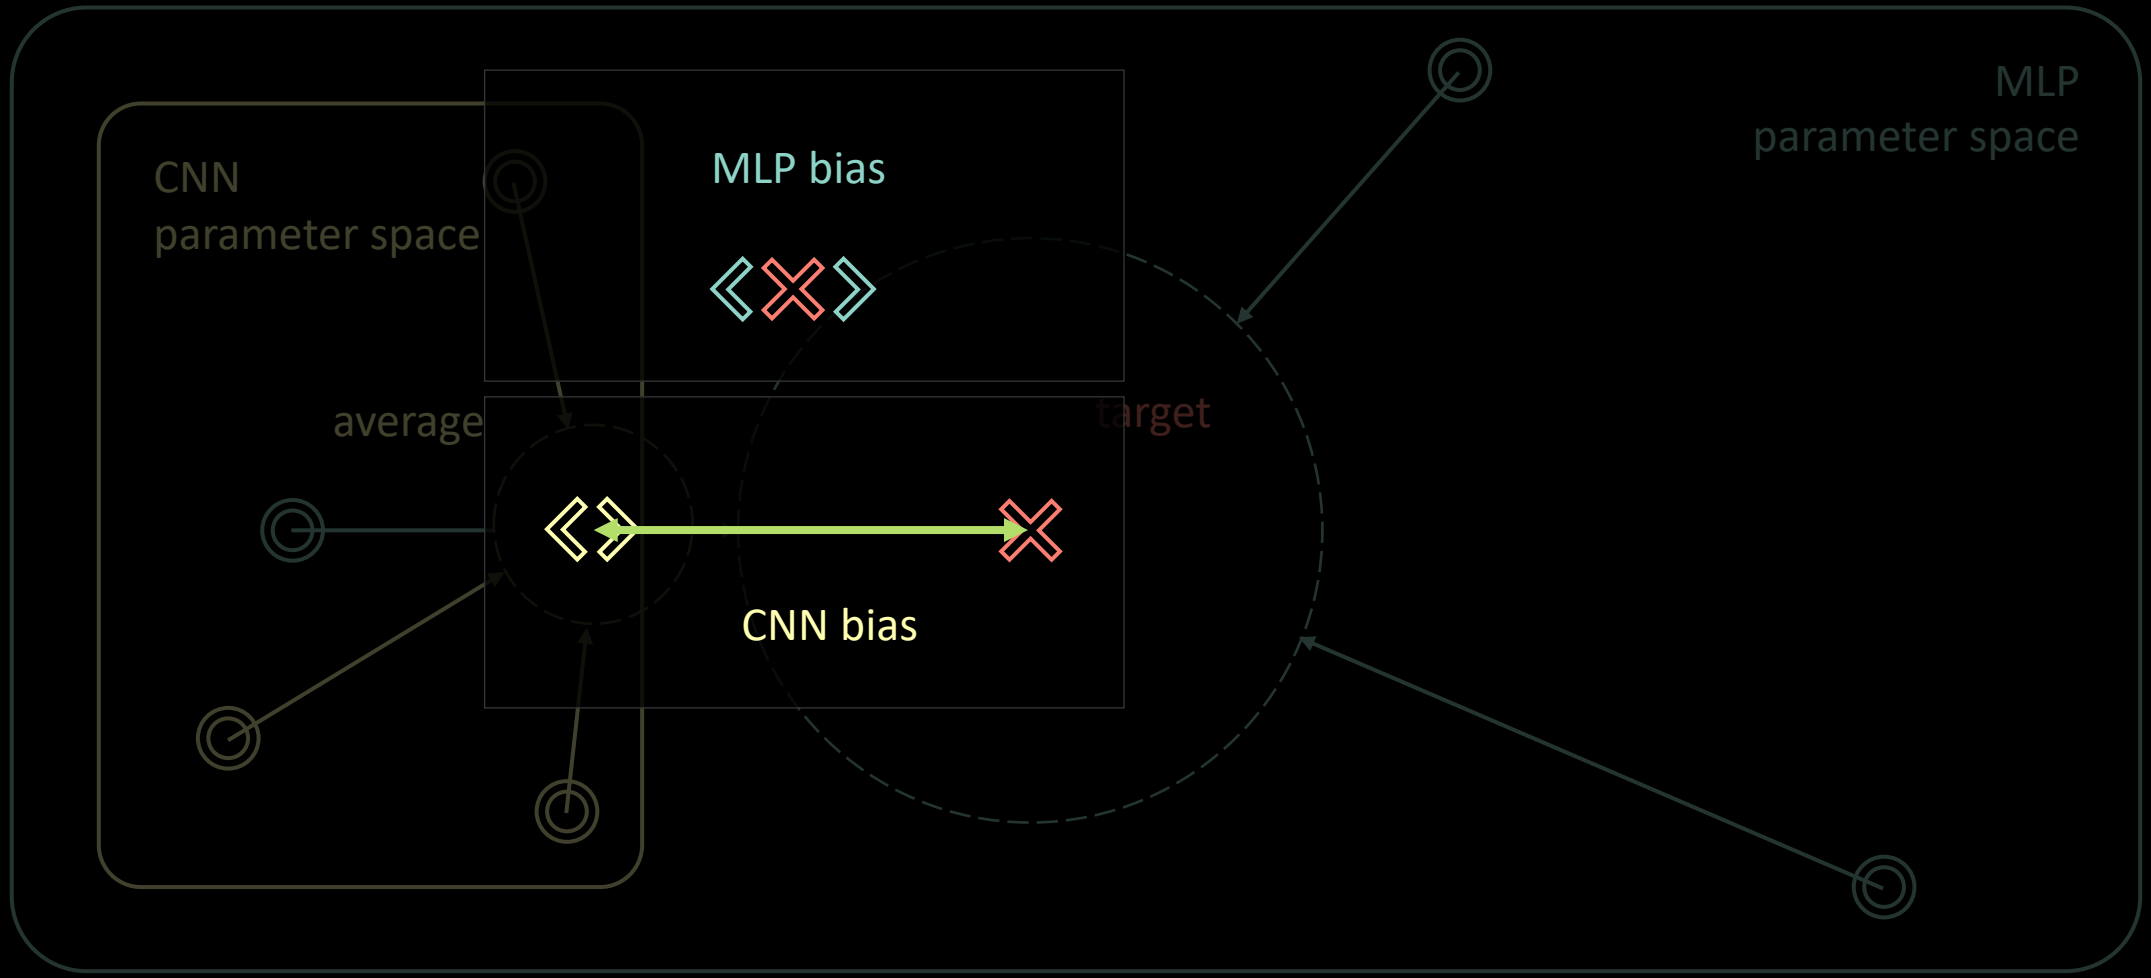

Kernels – 1D data

kernel size: $2 \times 7 \times 3$

1D data uses 3D kernels-collection!

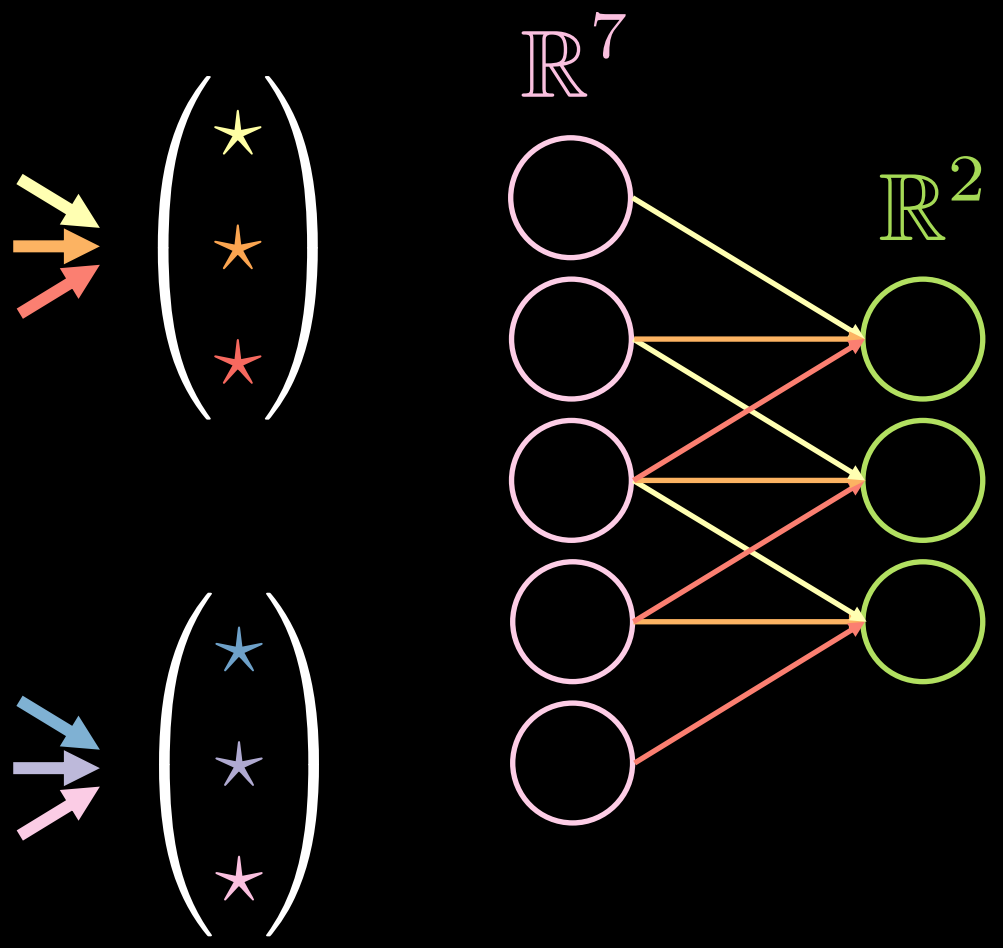

Padding – 1D data

kernel size: $2 \times 7 \times \boxed{3}$

zero padding: $(\boxed{3} - 1)/2 = \boxed{1}$

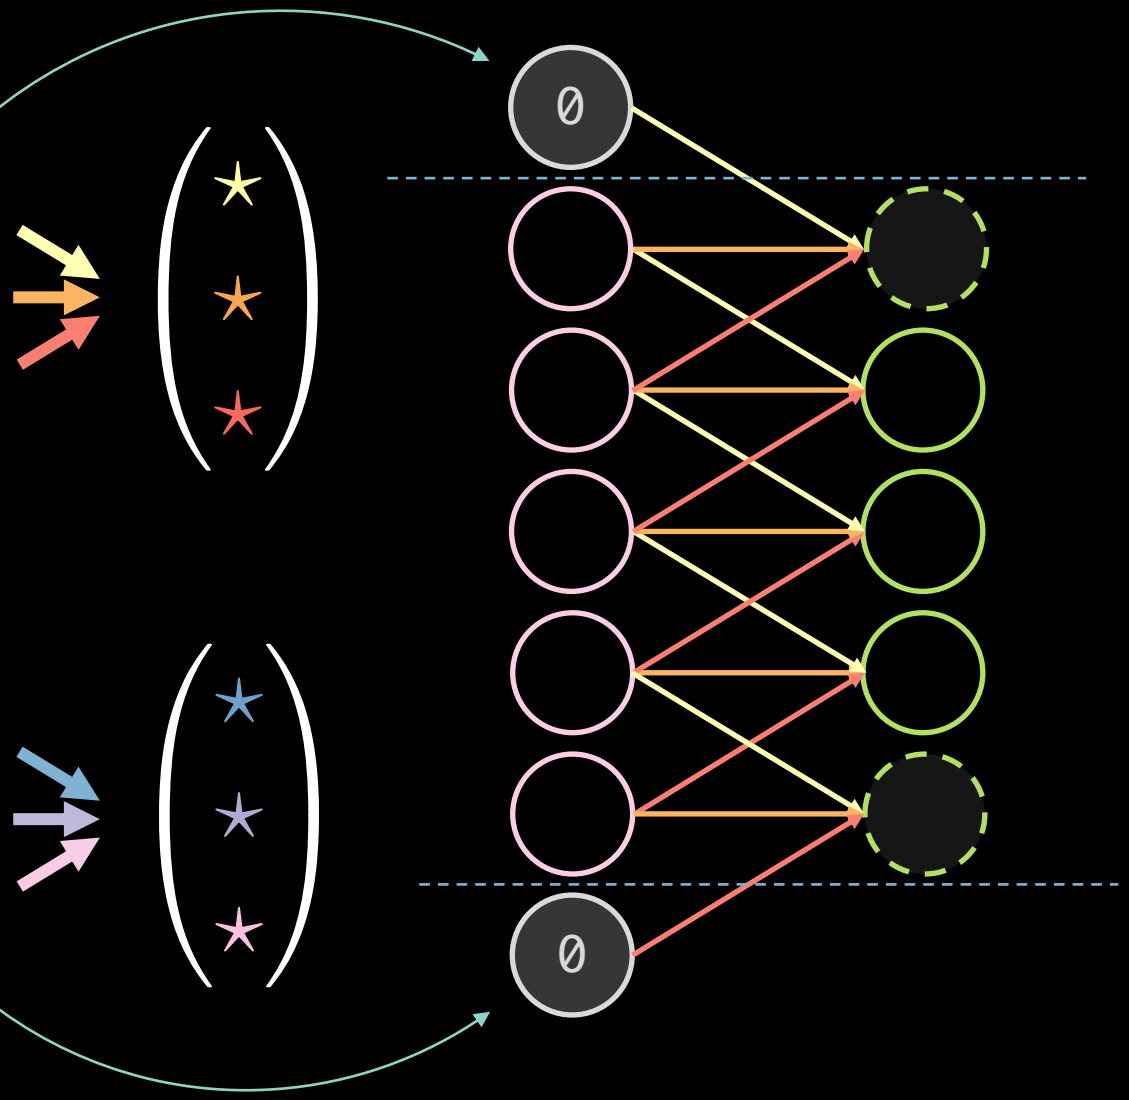

Standard spatial CNN

- Multiple layers
  - Convolution
  - Non-linearity (ReLU and Leaky)
  - Pooling
  - Batch normalisation
- Residual bypass connection

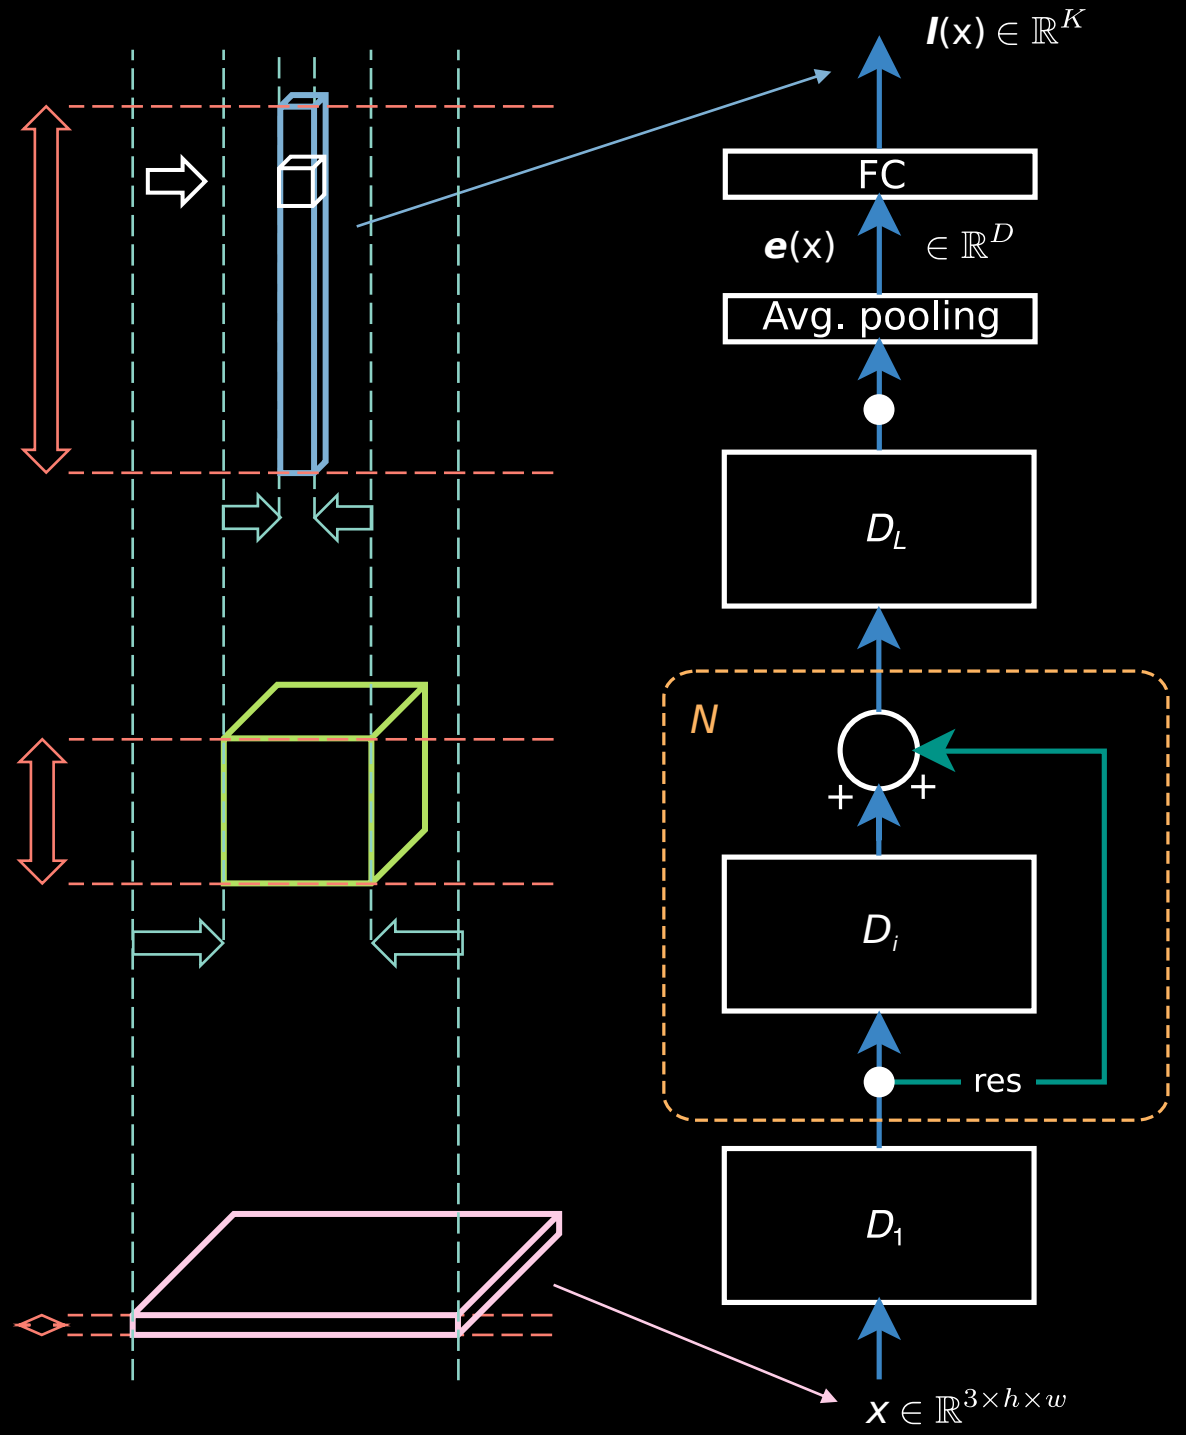

Pooling

$$\|\underline{x}\|_p := \left(\sum_i |x_i|^p \right)^{1/p}$$

$$\|\underline{x}\|_p \to \max(\underline{x}), \; p \to +\infty$$