Bayesian neural nets

Estimating a predictive distribution

Yarin Gal & Zoubin Ghahramani (2016)
Dropout as a bayesian approximation: Representing model uncertainty in deep learning

Why to care about uncertainty

- Cat / dog classifier, classify an hippopotamus
- Reliability in steering control
- Physics simulator prediction
- Minimising action randomness when connected to a reward

NN with dropout

$$\boldsymbol{h} = f(\boldsymbol{W}_h \boldsymbol{x} \odot \boldsymbol{\delta}_x + \boldsymbol{b}_h)$$

$$\hat{\boldsymbol{y}} = g(\boldsymbol{W}_y \boldsymbol{h} \odot \boldsymbol{\delta}_h + \boldsymbol{b}_y)$$

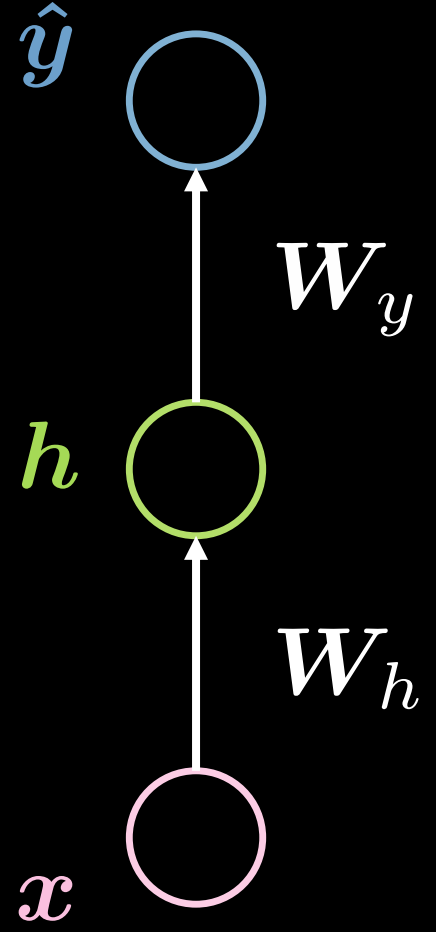

$$(\boldsymbol{\delta}_x)_i, (\boldsymbol{\delta}_h)_j \sim \mathrm{Ber}(1-r)$$

$$\boldsymbol{\delta} \leftarrow \frac{1}{1-r}\boldsymbol{\delta}$$

dropping rate

$$\boldsymbol{x}, \boldsymbol{\delta}_x \in \mathbb{R}^n$$

$$\boldsymbol{h}, \boldsymbol{\delta}_h \in \mathbb{R}^d$$

$$\hat{\boldsymbol{y}} \in \mathbb{R}^K$$

$$\boldsymbol{W}_h \in \mathbb{R}^{d \times n}$$

$$\boldsymbol{W}_y \in \mathbb{R}^{K \times d}$$

$$f(\cdot), g(\cdot) : (\cdot)^+, \sigma(\cdot), \tanh(\cdot), \mathrm{softmax}(\cdot)$$

Regression (I)

- See notebook demo

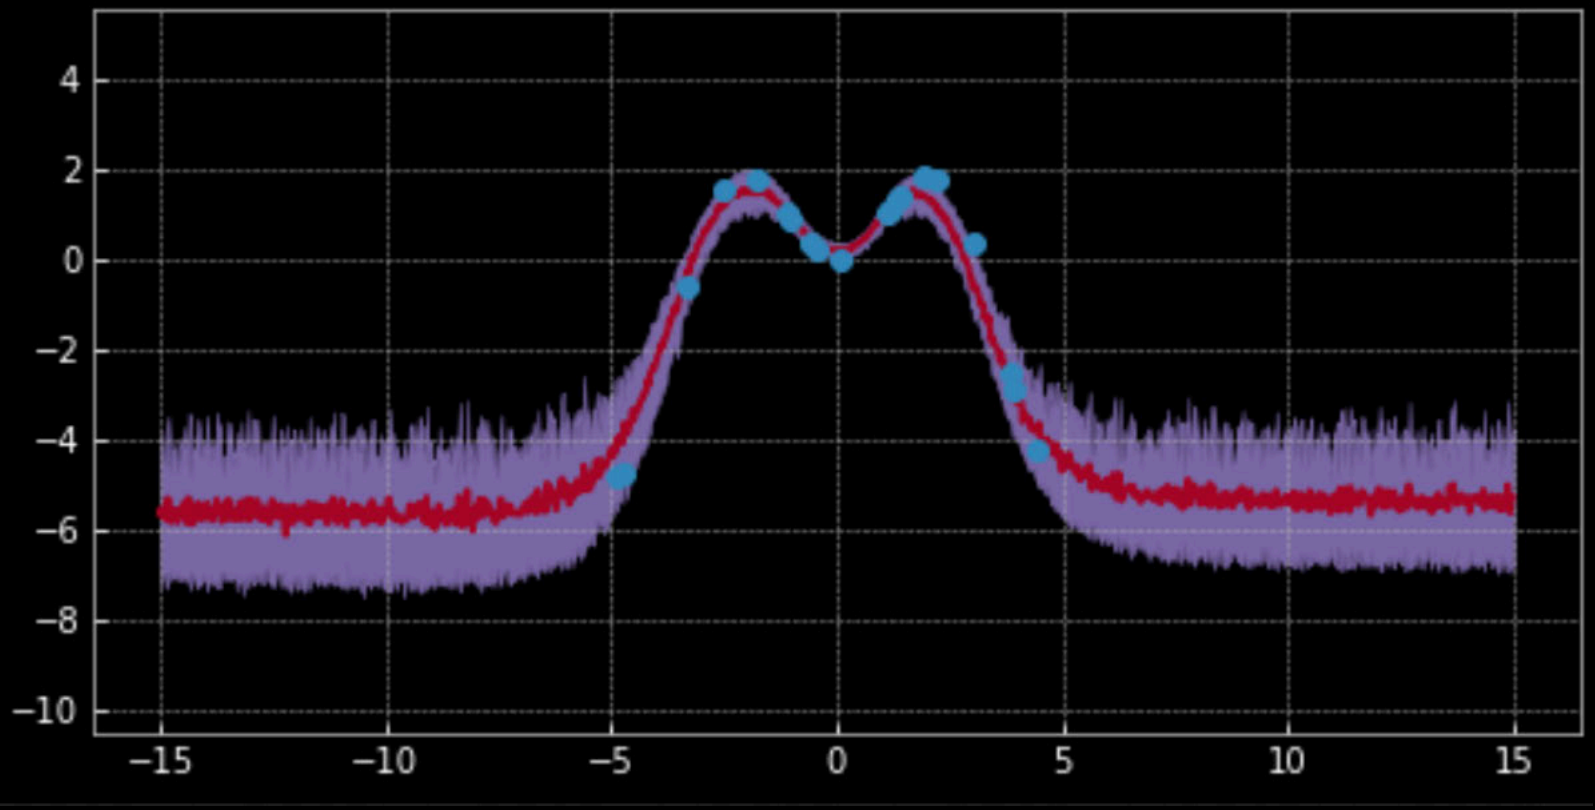

Regression (II)

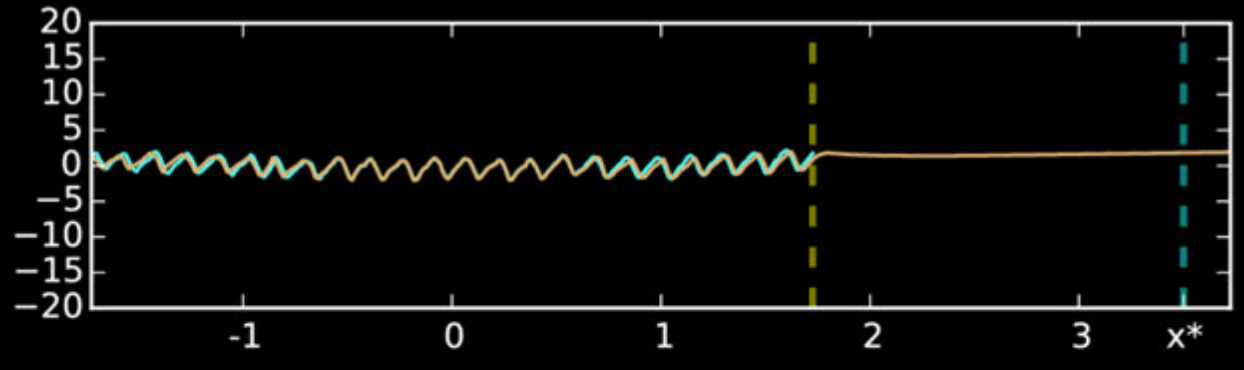

(a) Standard dropout with weight averaging

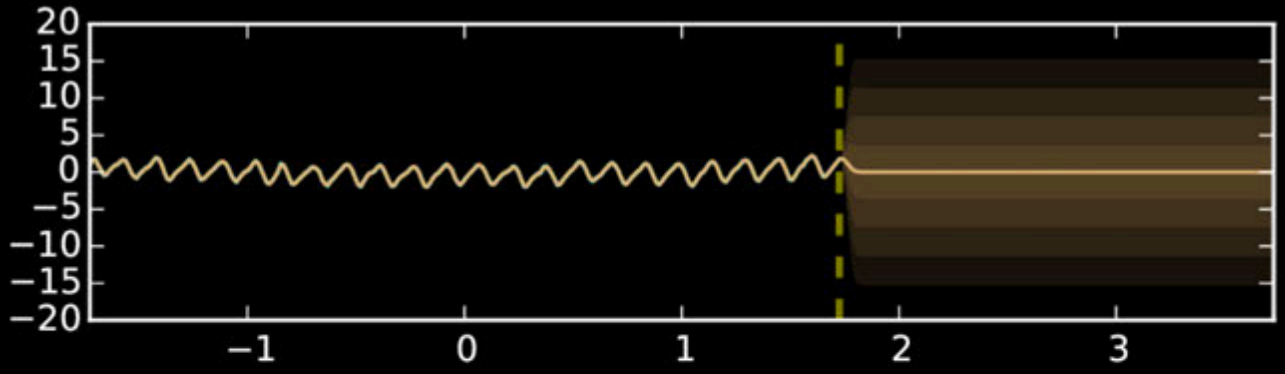

(b) Gaussian process with SE covariance function

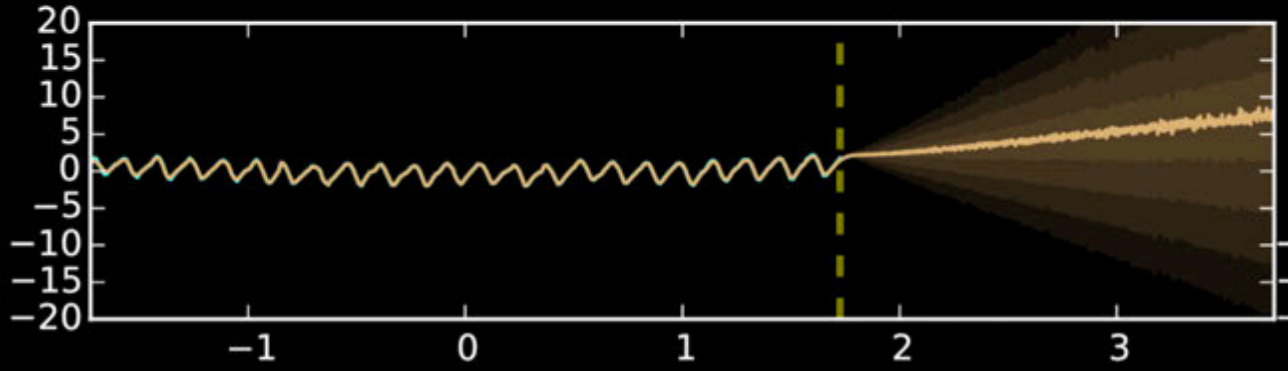

(c) MC dropout with ReLU non-linearities

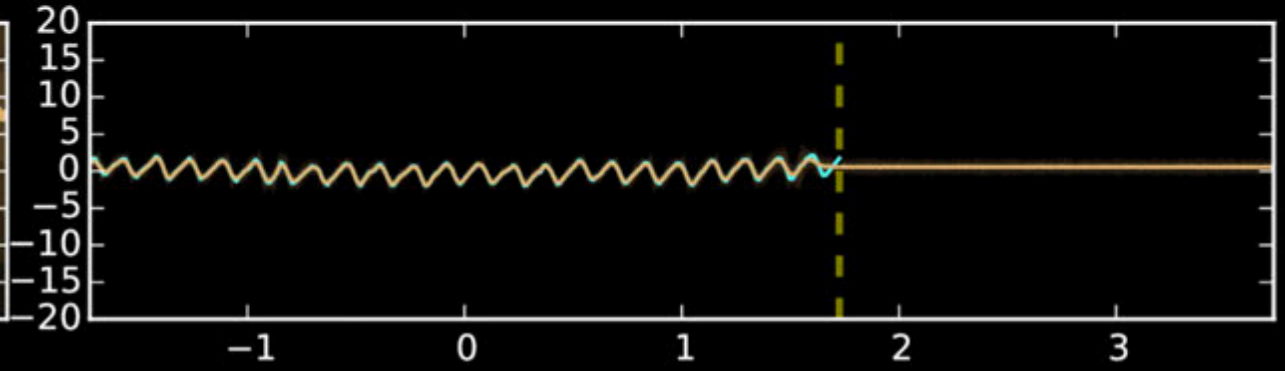

(d) MC dropout with TanH non-linearities

Binary classification

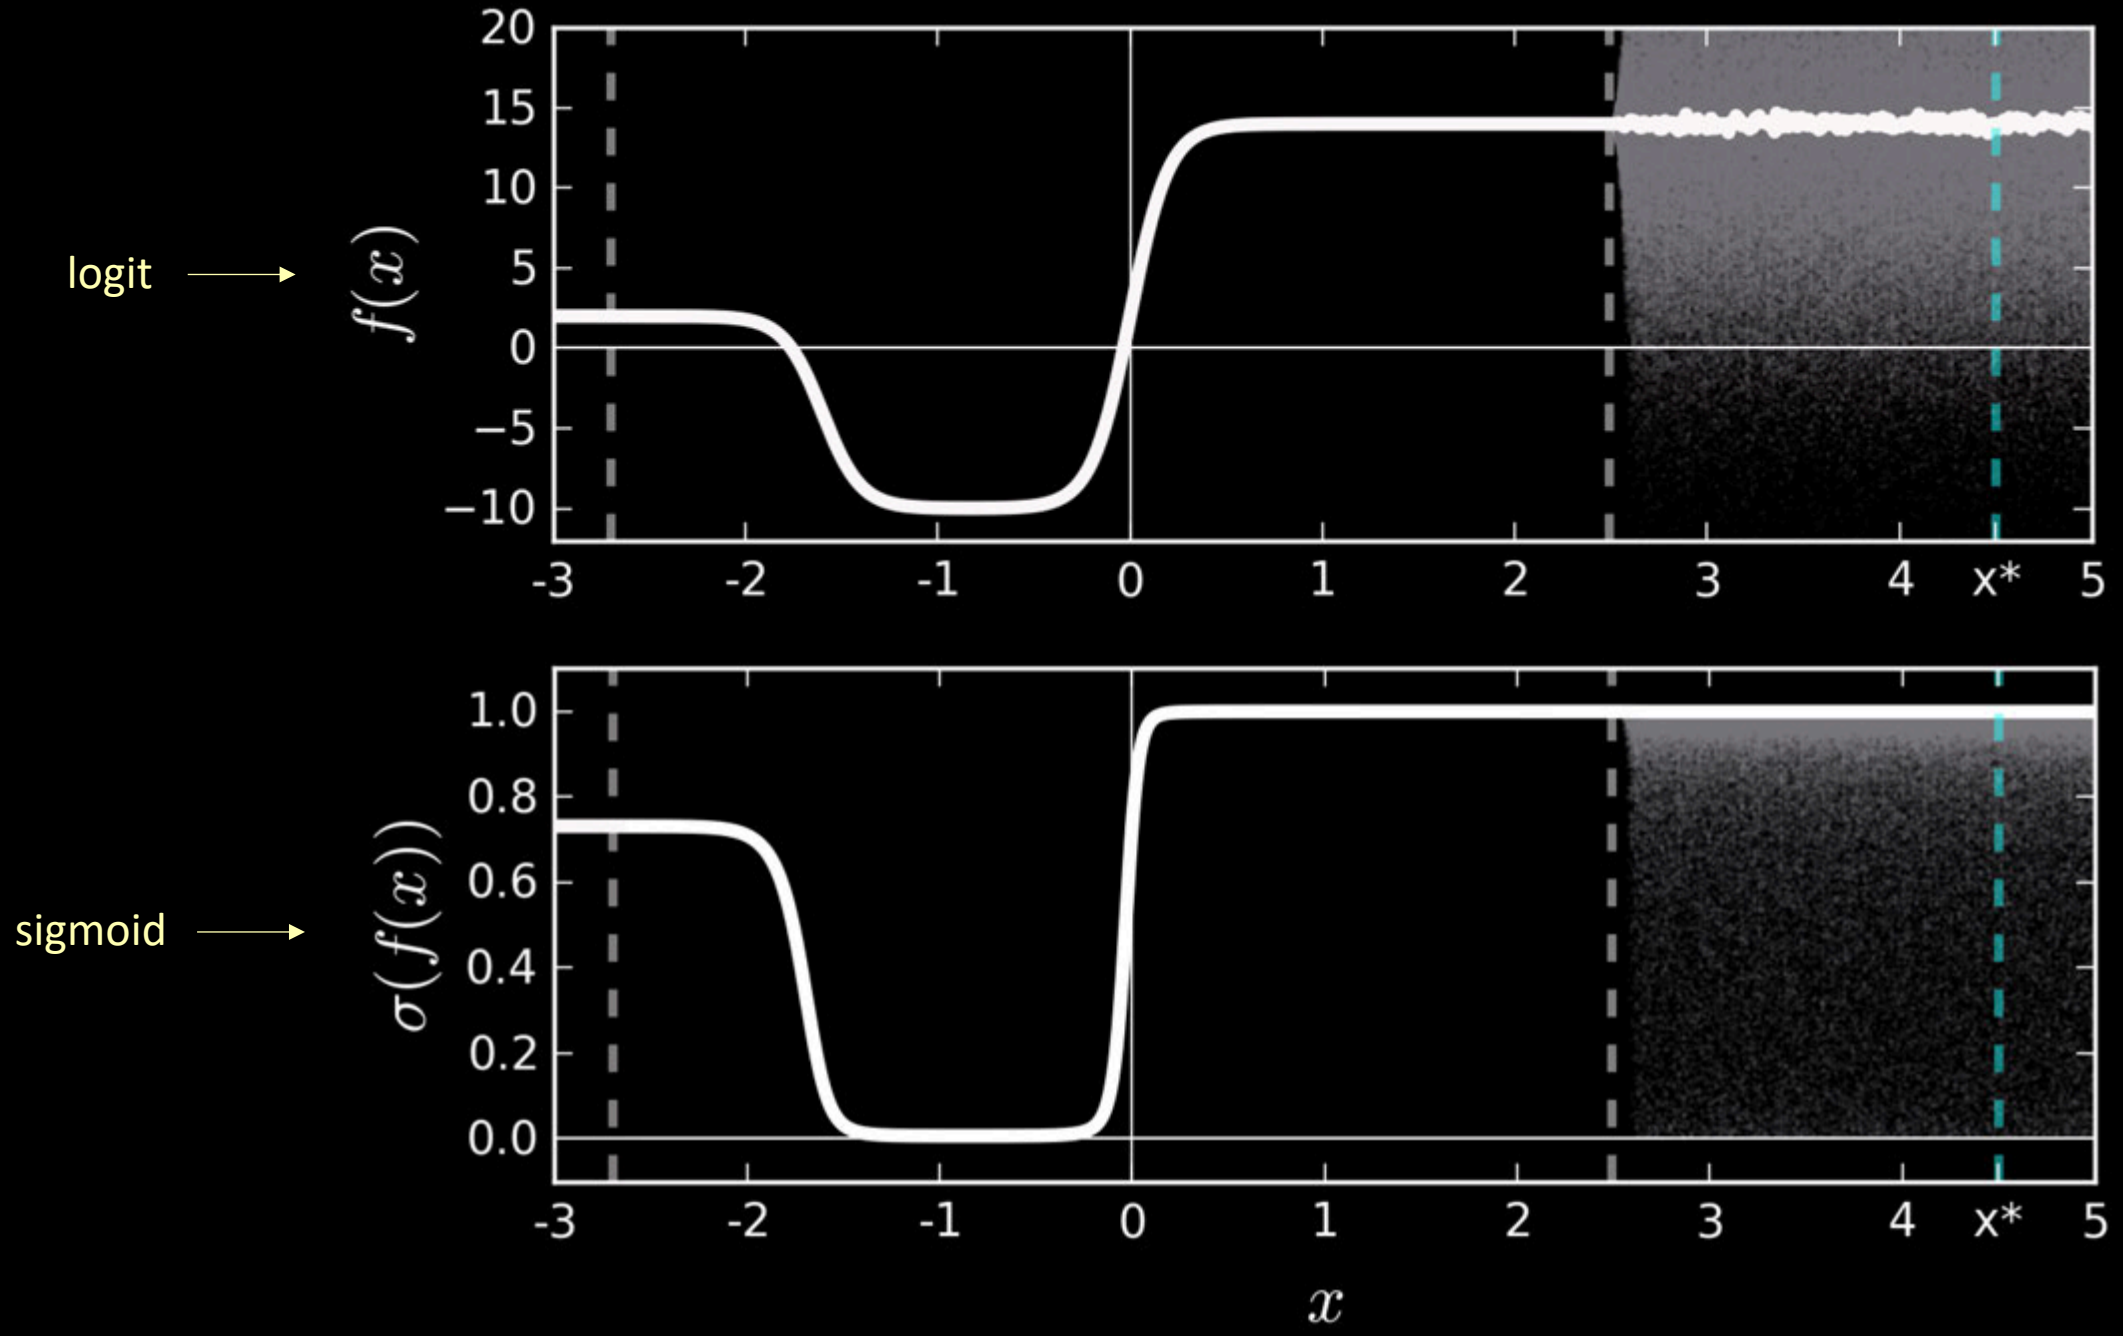

Multi-class classification